



7 Beyond Row Stores



Column-Stores

- Are also based on the relational data model
 - use of relational algebra (and their operators) as a foundation
 - Column stores are (just) an alternative physical data model
- Follows a logical query plan similar to row-stores model
- Column-based compression allows for data processing without decompression
 - Lossless compression techniques (e.g. Lempel-Ziv and derivatives) → identical (uncompressed) values imply identical compressed representation, i.e. value within search predicate is compressed beforehand and subsequently compared with the compressed representation
 - Order-preserving compression schemes (like RLE) → search values are directly comparable or similar mechanism to lossless-compression, i.e. aggregate functions like MIN/MAX or SUM can be processed in a compressed format.

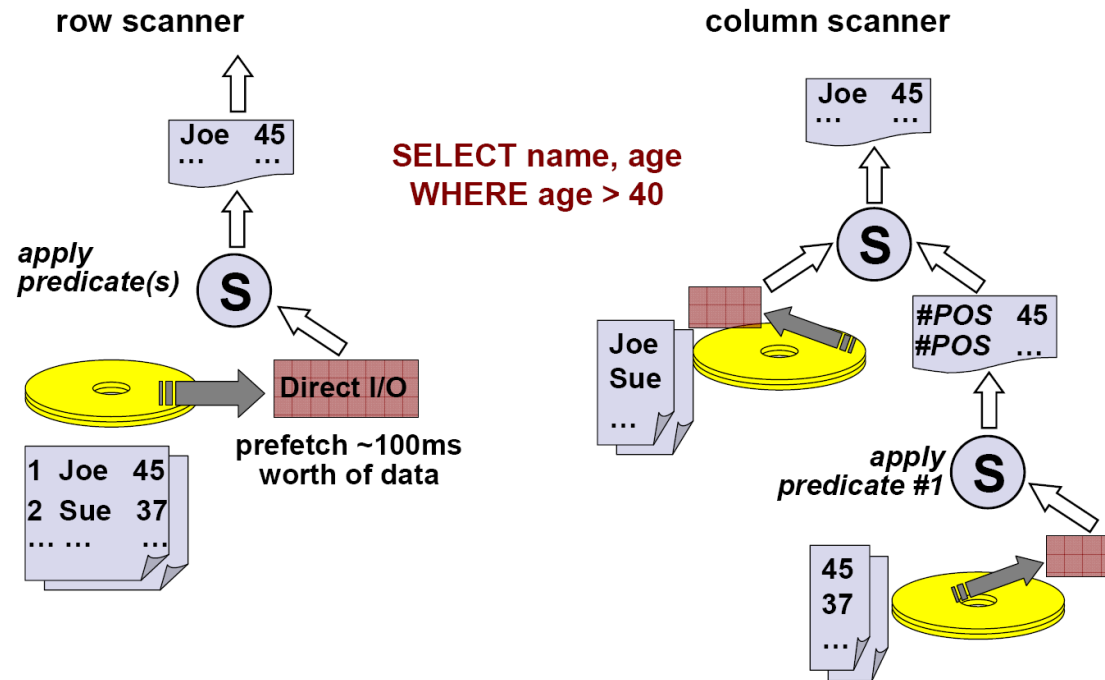


Column scan operators

- Translate value position information into disk locations
- Combine and reconstruct (when needed) partial or entire tuples out of different columns (*Materialization*)

Join operators

- Can either rely on column-scanners for receiving reconstructed tuples, or they can operate directly on columns by first computing a join index and then fetching qualifying value



- Reads from a single file
- Iterates over pages, for each page iterates over tuples and applies predicates

- Must read as many files as are columns specified in the query
- Series of pipelined scan nodes
- Applies predicates by reading a column, creating {position, value} pairs for all qualified tuples
- Attaches values corresponding to input positions from other columns



```
SELECT AVG (Attr. 1), Attr. 2 FROM Tabelle
WHERE Attr. 2=X;
```

Row Store

Seite

111	104	121	...		T1:
...	T2:	Attr. 1	Attr. 2 = X	...	T3:	Attr. 1	...
Attr. 2	...	T4:	Attr. 1	Attr. 2 = X	...	T5:	Attr.
1	Attr. 2	...					

Selektion Attr. 2

T2:	Attr. 1	Attr. 2 = X	...
T4:	Attr. 1	Attr. 2 = X	...

AVG(Attr. 1)
+ Projektion

Ergebnis

Column Store

Seite

211	199	222	...	
	T1:	Attr. 1	T2:	Attr.
1	T3:	Attr. 1	T4:	...

AVG(Attr. 1)

AVG(Attr. 1)

Seite

45	32	101	...
T1:	Attr. 2		T2:
Attr. 2 = X		T3:	Attr. 2
	T4:	Attr. 2 = X	

Selektion

T2:	Attr. 2 = X
T4:	Attr. 2 = X

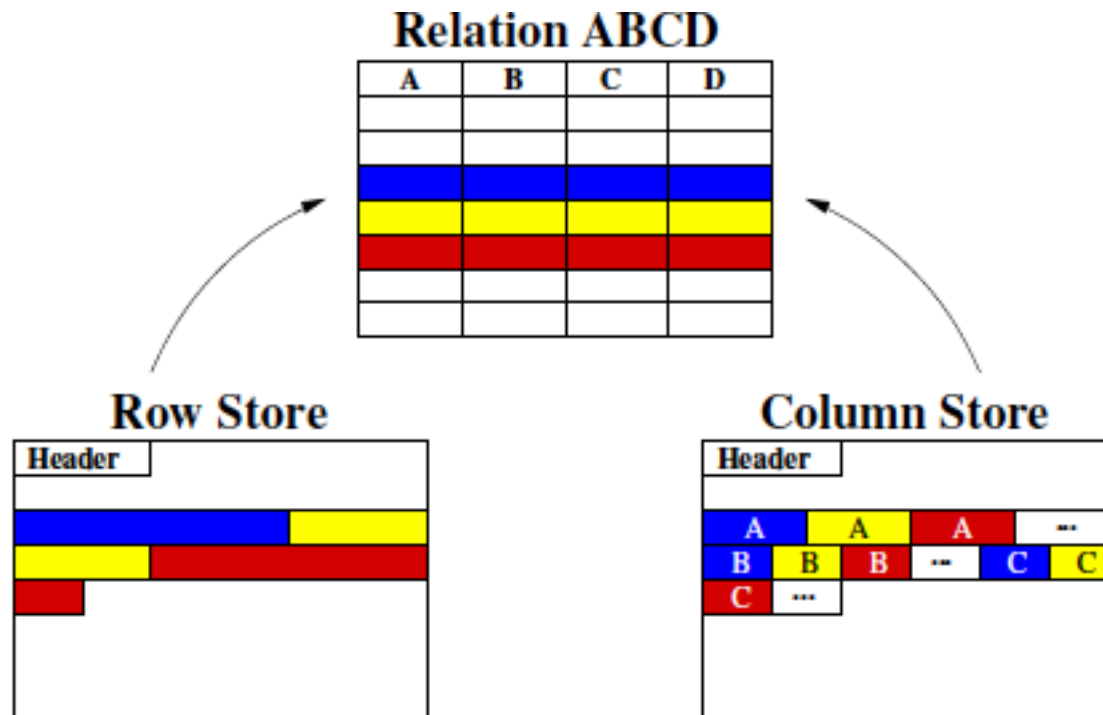
Merge Columns

Ergebnis



Operator

- Heißt SPC (Scan, Predicate, Construct) für volle Tupelrekonstruktion
- Merge ist ein k-tuplige Rekonstruktion (mit Spalten VAL1, ..., VALk)





Frühe Materialisierung (EM)

- Anfrageverarbeitung sehr nah an Row-Stores
- Aggregatfunktionen auf einzelnen Columns
- Tupelrekonstruktion sobald Tupel verwendet
- Zumeist Verwendung bei tupel-orientierter Anfragebearbeitung

Späte Materialisierung (LM)

- So lang wie möglich auf Columns arbeiten
- Mehrfacher Zugriff auf Basistabellen und/oder Zwischenergebnisse
- Folge: Anfrageplan kein Baum mehr
- Aber: Gleichzeitig Bearbeitung auf komprimierten und unkomprimierten Daten möglich
- Notwendig für effektive spalten-orientierte Anfragebearbeitung



Query Processing - Tuple Materialization

> When should columns be projected?



Row-store

- Column projection involves removing unneeded columns from tuples
- Generally done as early as possible

Column-store

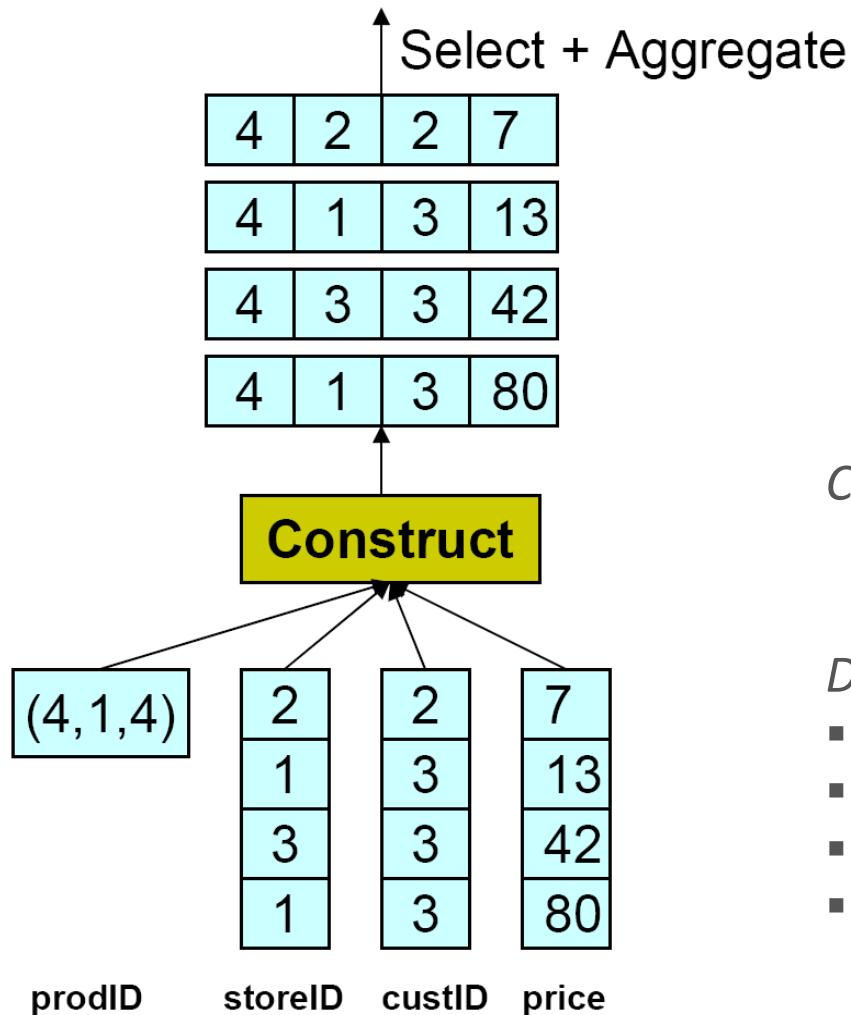
- Operation is almost completely opposite from a row-store
- Column projection involves reading needed columns from storage and extracting values for a listed set of tuples (Materialization)

Early materialization

- Project columns at beginning of query plan
- Straightforward since there is a one-to-one mapping across columns

Late materialization

- Wait as long as possible for projecting columns
- More complicated since selection and join operators on one column obfuscates mapping to other columns from same table



QUERY:

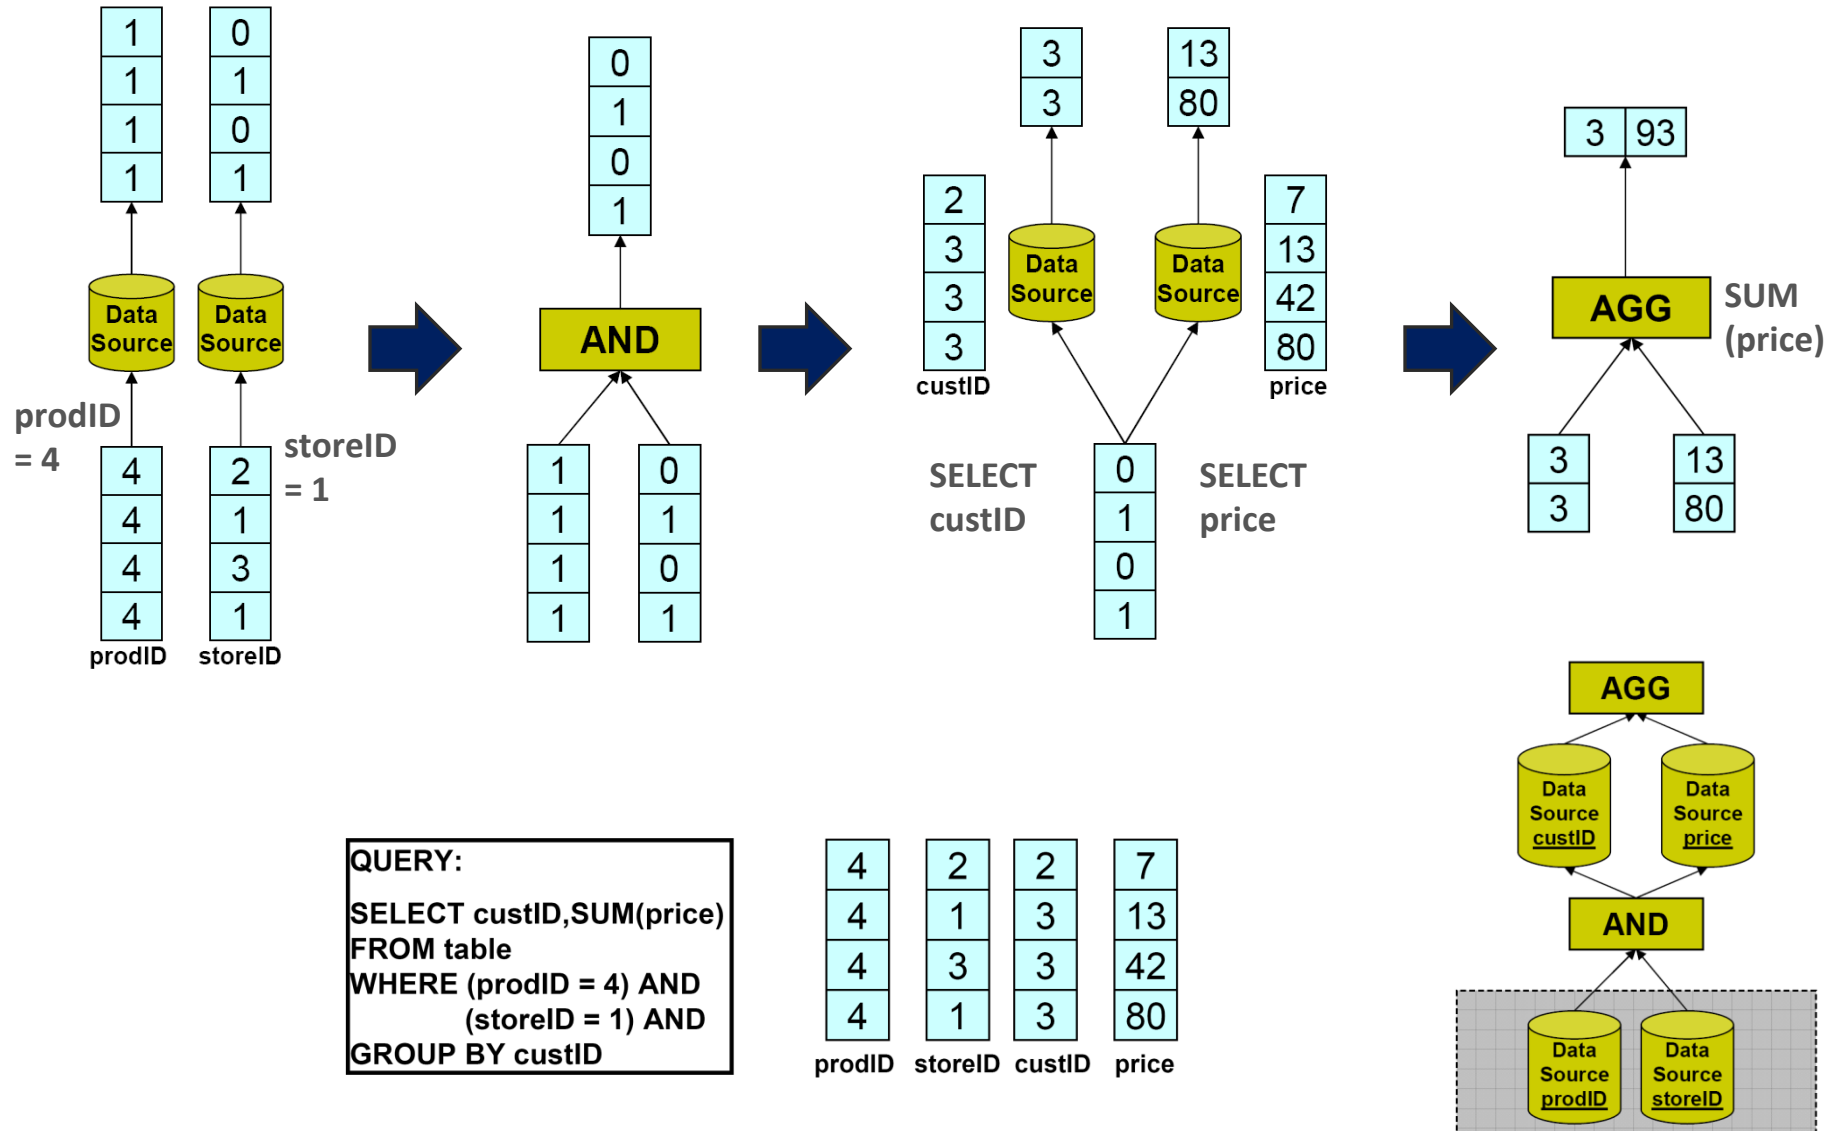
```
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1) AND
GROUP BY custID
```

Create rows first

Disadvantages:

- Need to construct all tuples
- Need to decompress data
- Poor memory bandwidth utilization
- Loose opportunity for vectorized operation

> Late Materialization Example





Operate directly on columns

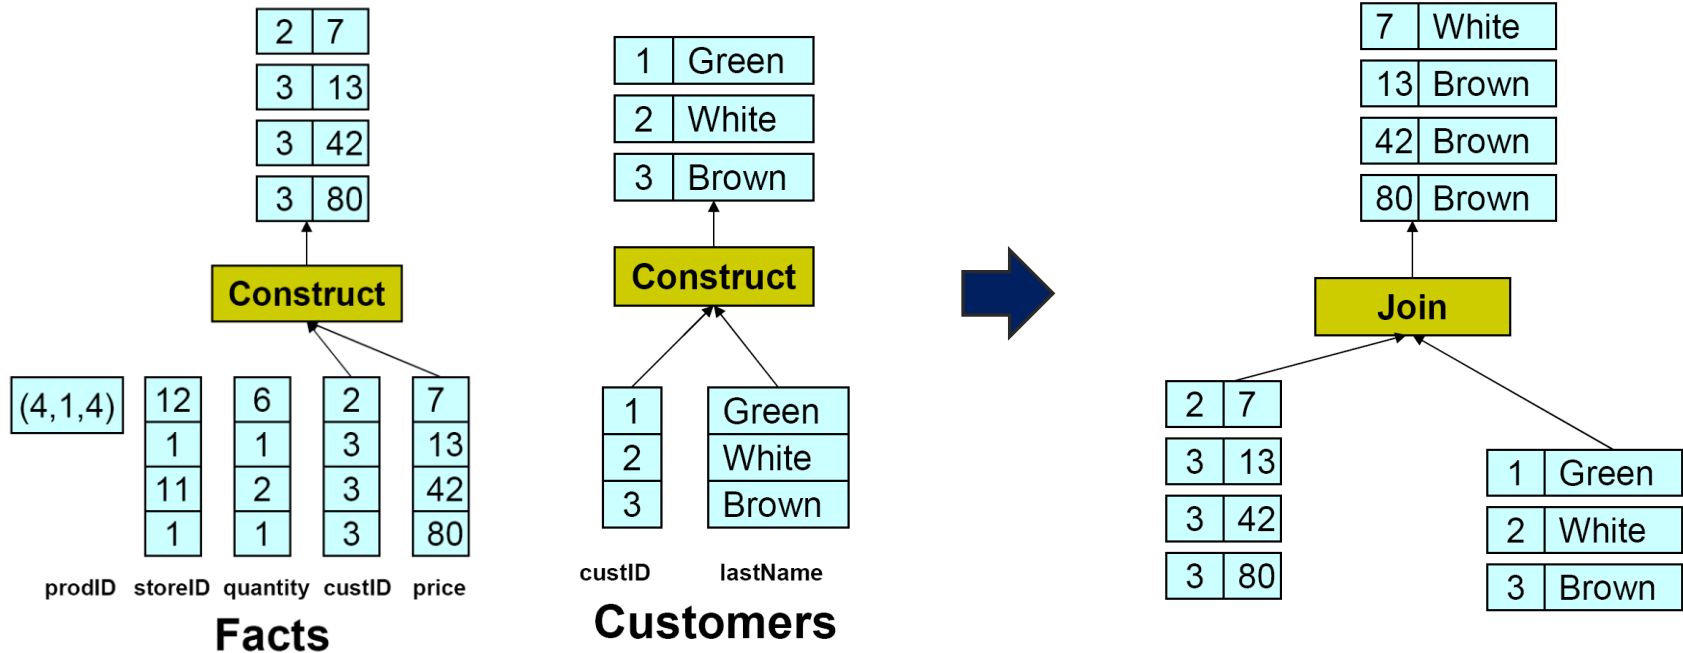
- Intermediate “position” lists need to be constructed in order to match up operations that have been performed on different columns

Advantages

- *Construct only relevant tuples, avoid unnecessary tuple construction*
- *Column can be kept compressed in memory*
 - Operating directly on compressed columns possible
- *Looping through column-oriented data tends to be faster than looping through tuples*
 - Values of the same column fill an entire cache line
 - Vector processing for column block access

Disadvantage

- *Columns may need to be accessed multiple times in a query plan*
 - Trade-off between late materialization optimizations and column reaccess costs



QUERY:

```
SELECT C.lastName, SUM(F.price)
FROM facts AS F, customers AS C
WHERE F.custID = C.custID
GROUP BY C.lastName
```

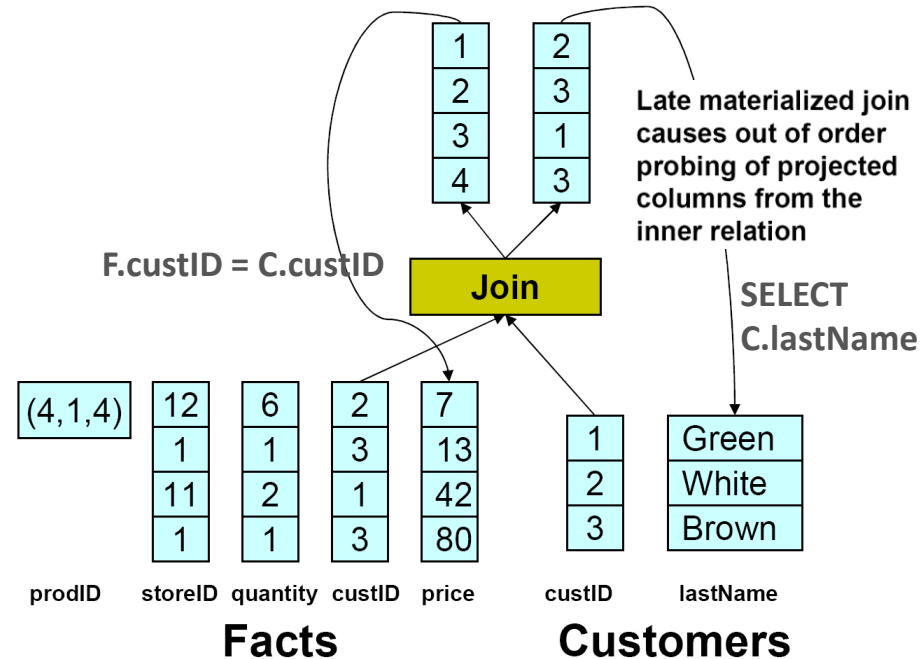
- Tuples have already been constructed before reaching the join operator
- Join functions as it would in a standard row-store system and outputs tuples



Results in two sets of positions, one for the fact table and one for the dimension table

42		
36		
42	38	
44	42	1 2
38	46	3 2
		5 1

- Indicates which pairs of tuples passed the join predicate
- At most one position list is produced in sorted order (fact table)
 - Merge join of positions can be used to extract other column values
- Values from dimension table need to be extracted in out-of-position order
 - Can be significantly more expensive



QUERY:

```
SELECT C.lastName,SUM(F.price)
FROM facts AS F, customers AS C
WHERE F.custID = C.custID
GROUP BY C.lastName
```



Invisible Join



Properties

- Late materialized join, minimizes the values that need to be extracted out-of-order
- Makes sure that the table that can be accessed in position order is the fact table for each join
- Rewrites joins into predicates on the foreign key columns in the fact table
- Position lists from the fact table are then intersected (in-position order)
- Reduces the amount of data that must be accessed out of order from the dimension tables



Phase 1

- For each predicate dimension table keys are extracted which satisfy the predicate
- Keys are used to build a hash table

Apply “region = ‘Asia’” On Customer Table

custkey	region	nation	...
1	ASIA	CHINA	...
2	EUROPE	FRANCE	...
3	ASIA	INDIA	...



Hash Table Containing
Keys 1 and 3

Apply “region = ‘Asia’” On Supplier Table

suppkey	region	nation	...
1	ASIA	RUSSIA	...
2	EUROPE	SPAIN	...



Hash Table Containing
Key 1

Apply “year in [1992,1997]” On Date Table

dateid	year	...
01011997	1997	...
01021997	1997	...
01031997	1997	...

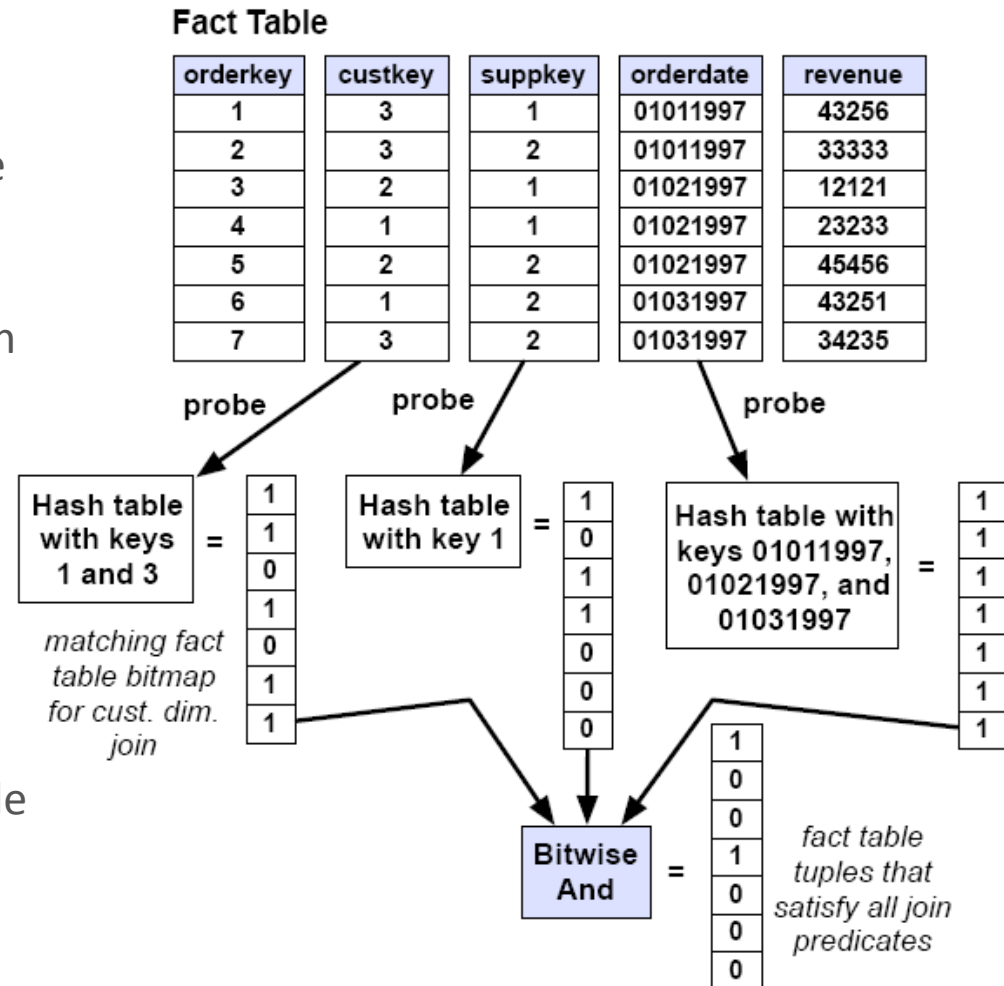


Hash Table Containing
Keys 01011997, 01021997,
and 01031997



Phase 2

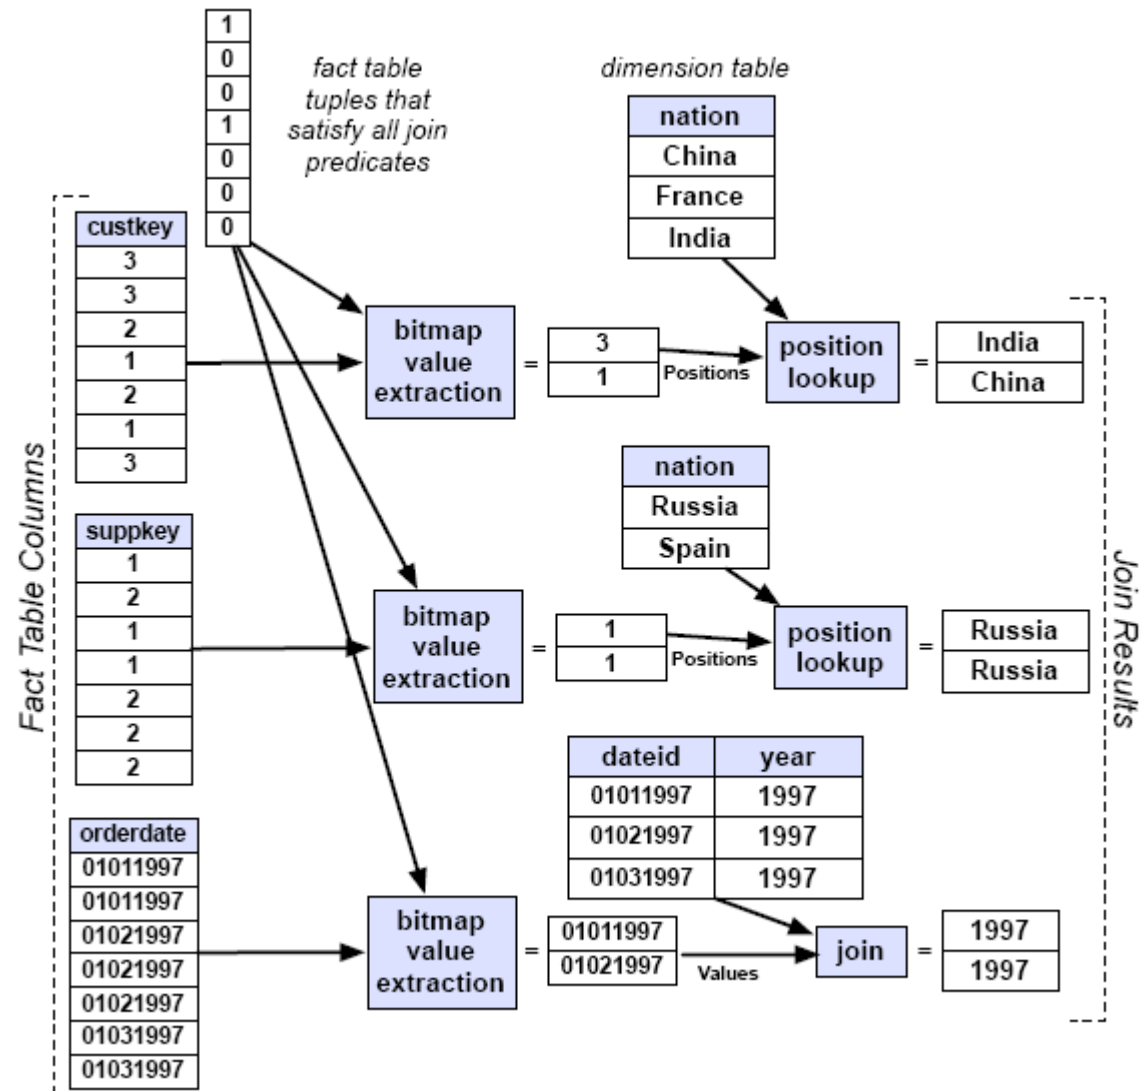
- Hash table is used to extract the positions of records in the fact table that satisfy the corresponding predicate
- Each value in the foreign key column of the fact table is probed into the hash table
- Results in a list of all positions in the foreign key column that satisfy the predicate
- Lists from all of the predicates are intersected to generate a list of satisfying positions P in the fact table





Phase 3:

- The third phase uses the list of satisfying positions P in the fact table to get foreign key values and hence needed data values from the corresponding dimension table.





Which Column-Specific Optimization is most significant?

> Most significant?



Late materialization

- Improves performance by a factor of 3.

Block iteration

- Improves performance by 5% - 50%.


Column-specific compression

- Improves performance by a factor of 2 on average
- Improves performance by a factor of 10 on queries that access sorted data.

Invisible join on star schema

- Improves performance by 50% - 70%

> Summary



- + No tuple reconstruction
- Massive virtual function calls
- High materialization costs


tuple(s)-at-a-time

Row-Stores

→ *Update friendly*

Column-Stores

→ *Less update friendly*



- + Low materialization costs
- High tuple reconstruction costs
- Single “next” call

column-at-a-time



Vectorized Execution

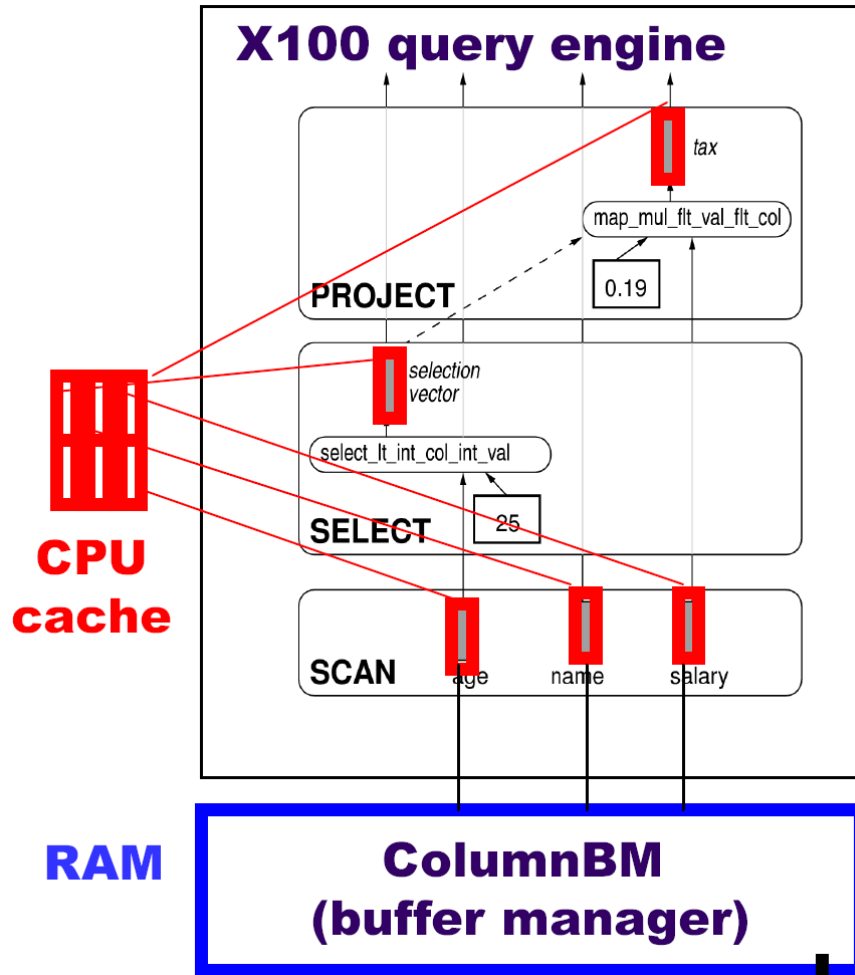


Block Iteration

- Blocks of values from the same column are sent to an operator in a single function call
- If the column is fixed-width, these values can be iterated through directly as an array
- Advantages
 - Minimizes per-tuple overhead
 - Exploits potential for parallelism on modern CPUs, e.g. loop-pipelining techniques

Example: MonetDB/X100

- New hyper-pipelining query execution engine for MonetDB
- In-cache vectorized processing
- MonetDB materializes all intermediate results -> limits scalability
- Volcano iterator pipelining model: used in most row-stores
 - But: tuple-at-a-time processing
- MonetDB/X100 combines benefits of low-overhead column-wise query execution with the absence of intermediate result materialization in the Volcano iterator model



Combines Volcano model with vector processing

All vectors together should fit the CPU cache

Optimizer tune vector size, given the query characteristics.

Vectors are compressed

- Decompression between RAM and CPU cache



In-cache execution

- Only „randomly“ accessible memory is the CPU cache
- Main memory only used in the buffer manager for buffering I/O operations and large intermediate results

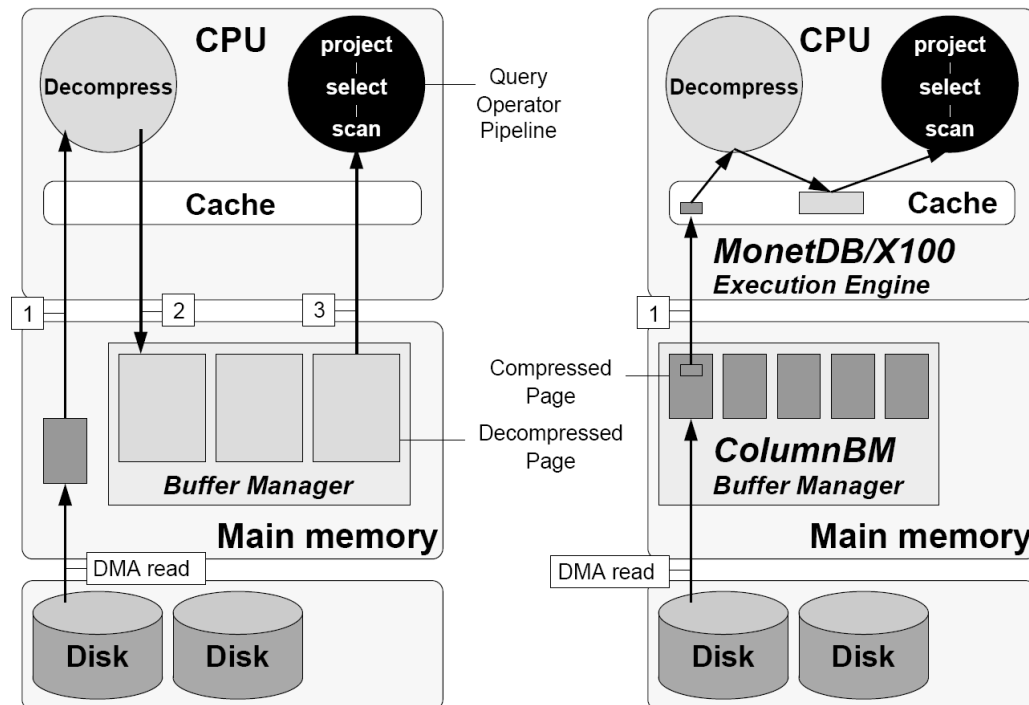
Vectorized execution

- Vertically decomposed tables are further partitioned horizontally into small chunks called *vectors*
- A set of aligned vectors (one for each attribute), representing a set of tuples, is a single data unit in the execution pipeline
- An optional selection vector contains the positions of tuples currently taking part in processing
- The control logic of operators is common for all data types, arithmetic functions, predicates etc.
- The actual data processing in the operators is performed by a set of *execution primitives* – simple, specialized and CPU-efficient functions



RAM-CPU Cache Compression

- Compression/decompression is used on the boundary between the CPU cache and RAM storage levels (instead of disk and RAM)
- Cache pages in the buffer manager (i.e. in RAM) in compressed form.
- Tuple values are decompressed at a small granularity (such that they fit the CPU cache) just-in-time, when the query processor needs them.



Advantages

- RAM-Cache boundary only crossed once
- More (compressed) data in RAM
- Less bandwidth use