



8 Adaptive Query Processing



Data independence facilitates modern DBMS technology

- Separates specification (“what”) from implementation (“how”)
- Optimizer maps declarative query → algebraic operations
- Platforms, conditions are constantly changing:
 $\Delta_{\text{app}} / \Delta t \ll \Delta_{\text{env}} / \Delta t$
- Query processing **adapts** implementation to runtime conditions
 - Static applications → dynamic environments


Traditional Optimization Is Breaking

- In traditional settings
 - Queries over many tables
 - Unreliability of traditional cost estimation
 - Success & maturity make problems more apparent, critical
- In new environments
 - e.g. data integration, web services, streams, P2P, sensor nets, hosting
 - Unknown and dynamic characteristics for data and runtime
 - Increasingly aggressive sharing of resources and computation
 - Interactivity in query processing



Example Query: $\sigma_{A=x \wedge B=y}(R) \bowtie S$


- $f(\sigma_{A=x}(R)) = 1/10$
- $f(\sigma_{B=y}(R)) = 1/10$
- $f(\sigma_{A=x \wedge B=y}(R)) = 1/10 \cdot 1/10 = 1/100$
(We assumed the predicates were independent by multiplying!)
- Problem 1: Correlation (functional dependency)
→ Problem: **underestimation** (leads to aggressive plans, e.g. NLJ)



Robustness of Plans

(Insensitivity to input statistics)

- **QEP1 (NLJ)**
 $C(\text{NLJ}) = |\sigma(R)| + |\sigma(R)| \cdot |S|$
- **QEP2 (SMJ)**
 $C(\text{SMJ}) = |\sigma(R)| \cdot \log_2 |\sigma(R)| + |S| \cdot \log_2 |S| + |\sigma(R)| + |S|$





Information Integration

- Query processing over heterogeneous systems and applications

Problem 1: Changing Workload Characteristics


- Unpredictable workload of external systems
- Infrastructural properties (network traffic, bandwidth)

Problem 2: Missing Knowledge about Statistics of External Systems

- Cardinalities, Selectivities, Ordering, Bandwidth
 - No access to statistical information
 - Estimation error increases exponentially in the size of the query
- ➔ Incremental statistic maintenance and re-optimization

All these problems lead to the same conclusion of **Adaptive Query Processing**

- Changing workload characteristics
- Unknown/uncertain statistics (external systems or correlation)

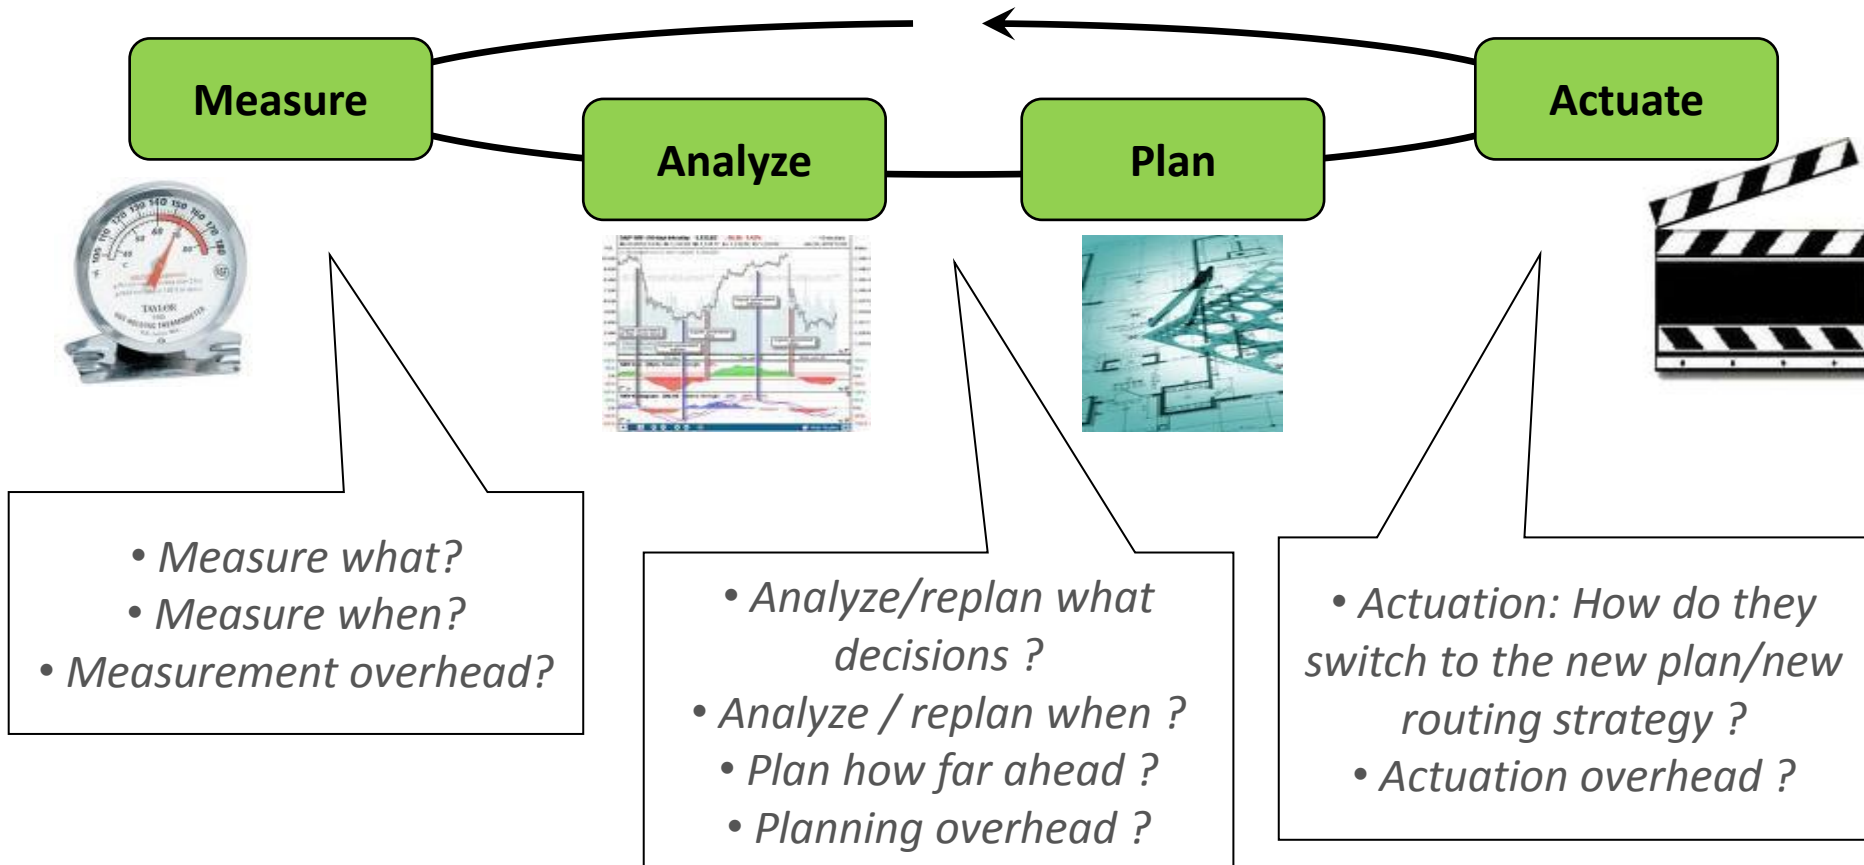




*MAPE (Monitor, Analyze, Plan, Execute) /
MAPA (Measure, Analyze, Plan, Actuate)*

[IBM. An architectural blueprint for autonomic computing. Technical report, IBM, 2005.]

[Zachary G. Ives, Amol Deshpande, Vijayshankar Raman: Adaptive query processing: Why, How, When, and What Next? VLDB 2007]





◆ CLASSIFICATIONS OF ADAPTIVE QUERY PROCESSING

- ❖ Inter-Query Adaptation (static, late binding)
- ❖ Inter-Operator Adaptation
- ❖ Intra-Operator Adaptation
- ❖ Tuple Routing
- ❖ Rewriting-Based Classification
- ❖ Reactive vs. Proactive Reoptimization
- ❖ Synchronous vs. Asynchronous Reoptimization

◆ ADAPTIVE QUERY PROCESSING (IN ACTION)

- ❖ Plan-Based Adaptation
- ❖ Continuous-Query-Based Adaptation
- ❖ Routing-Based Adaptation

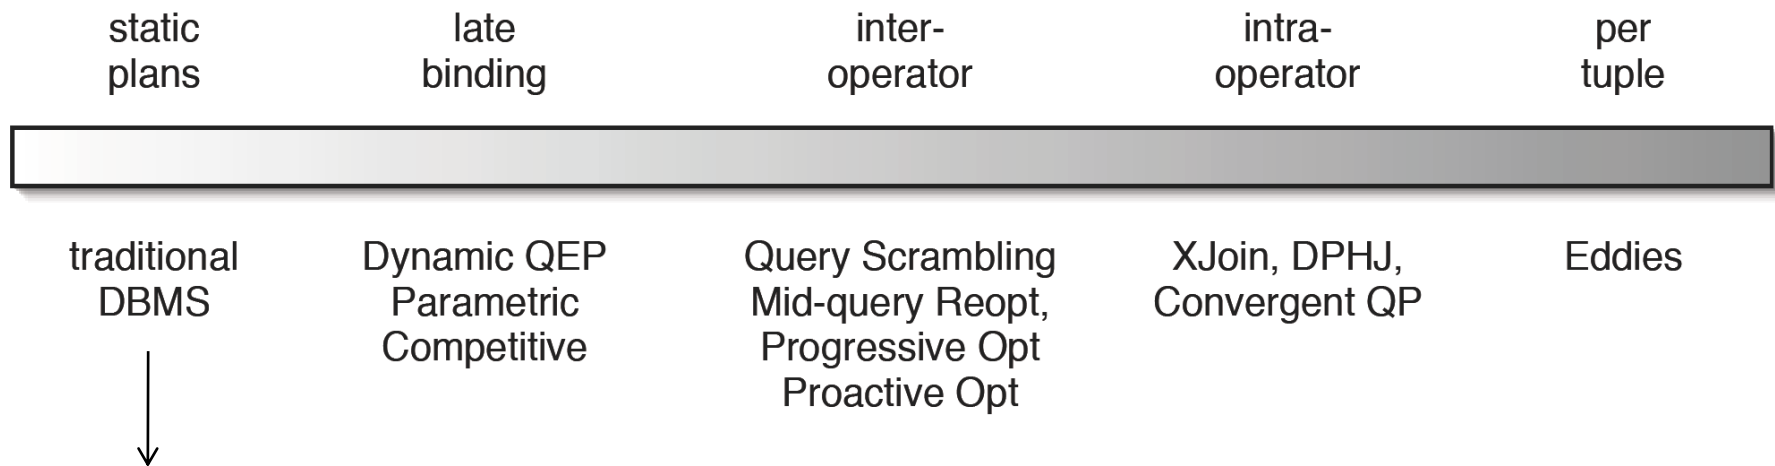


Classifications of Adaptive Query Processing



Spectrum of Adaptivity

[Amol Deshpande, Joseph M. Hellerstein, Vijayshankar Raman: Adaptive query processing: why, how, when, what next, SIGMOD 2006].



Inter-Query Optimization

- As established in System R

[Patricia G. Selinger, Morton M. Astrahan, Donald D. Chamberlin, Raymond A. Lorie, Thomas G. Price: Access Path Selection in a Relational Database Management System. SIGMOD 1979].


- UPDATE STATISTICS (cardinalities, histograms, index low/high keys)
- Dynamic Programming + Pruning Techniques

> Temporal Classification (1)



Example *Inter-Query* Adaptivity

- Opt for each query
- Opt for batch of queries



Query 1: \longrightarrow

SELECT *


FROM R,S,T,U

WHERE R.a=S.a AND

R.a=T.a AND

R.a=U.a

Full Optimization




Update Statistics: \longrightarrow

RUNSTATS ON TABLE R ON
ALL COLUMNS

RUNSTATS ON TABLE S ON ALL
COLUMNS

RUNSTATS ON TABLE T ON ALL
COLUMNS

RUNSTATS ON TABLE U ON
ALL COLUMNS



Query 2: \longrightarrow

SELECT *


FROM R,S,T,U

WHERE R.a=S.a AND

R.a=T.a AND

R.a=U.a

Full Optimization



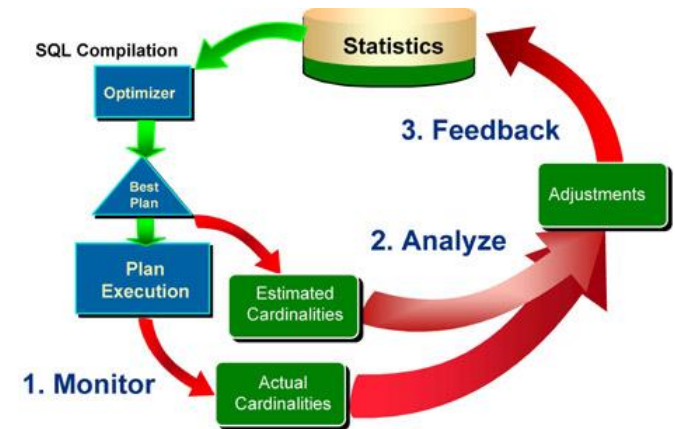


Inter-Query Adaptivity

LEO: DB2 Learning Optimizer

- E.g., correlation adjustments with maximum entropy approach

[Michael Stillger, Guy M. Lohman, Volker Markl, Mokhtar Kandil: LEO - DB2's LEarning Optimizer. VLDB 2001]



ASE: Adaptive Selectivity Estimation


[Chung-Min Chen, Nick Roussopoulos: Adaptive Selectivity Estimation Using Query Feedback. SIGMOD 1994]

- Real attribute value distribution approximated by curve-fitting function using a query feedback mechanism

SIT: Statistics on query expressions

[Nicolas Bruno, Surajit Chaudhuri: Exploiting statistics on query expressions for optimization. SIGMOD 2002]

- Exploit statistics build on expressions corresponding to intermediate nodes of a query plan rather than only base table statistics
- Problem of many relevant candidate statistics → workload-driven technique





Example *Late Binding*

- Use natural blocking points within a plan for reoptimization

Query:

SELECT *


FROM R,S,T,U

WHERE R.a=S.a AND





R.a=T.a AND

R.a=U.a

GROUP BY R.a



Re-Optimization of remaining subplan






Late Binding (Parametric Query Optimization)

Problem

- Unknown predicates at query compile-time (e.g., prepared statements)
- Similar to unknown or misestimated statistics
- Re-optimization for each incoming query may produce redundant work
- *Optimize-Once* and *Optimize-Always* both have disadvantages




Core Idea

- Optimize a query into a number of candidate plans (initial effort)
- Each candidate is optimal for some region of the parameter space
- Chooses an appropriate plan during query execution when the actual parameter values are known (can be much faster than re-optimization)


Approaches

- Progressive PQO (pay-as-you-go creation of candidates)
- PQO for linear cost functions
- AniPQO for non-linear cost functions

[Pedro Bizarro, Nicolas Bruno, David J. DeWitt: Progressive Parametric Query Optimization. IEEE Trans. Knowl. Data Eng., TKDE 2009]



> Temporal Classification (3)



Example *Inter-Operator*

- Insert artificial materialization points for reoptimization at arbitrary points between operators of a plan

Query:


SELECT *

FROM R,S,T,U

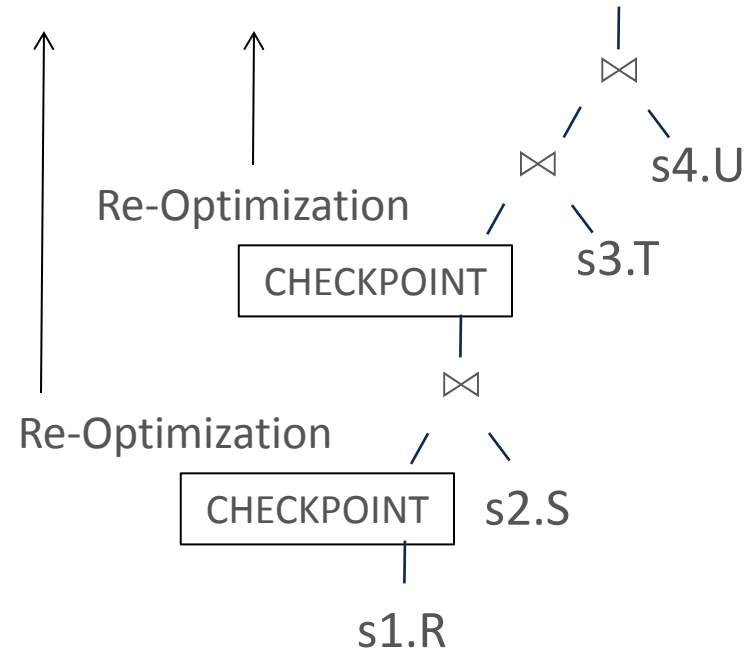
WHERE R.a=S.a AND

R.a=T.a AND


R.a=U.a



Optimal remaining subplans




> Temporal Classification (4)





Example *Intra-Operator*

- Use different plans for different partitions of the data
- Combine the results of subplans in stitch-up-phases

(a) an aggregation/join query as a combination of two plans plus a final stitch-up-phase



(b) a generalization:
„Pipelined Hash Join“ and
„Complementary Join Pair“



> Temporal Classification (5)



static plans late binding inter-operator intra-operator **per tuple**

Example *Tuple Routing*

- No plan (no fixed execution order of operators)
- Tuples are explicitly routed across relevant operators using routing policies

Query:

SELECT *




FROM R,S,T,U

WHERE $\sigma_p(T)$ AND

R.a=S.a AND

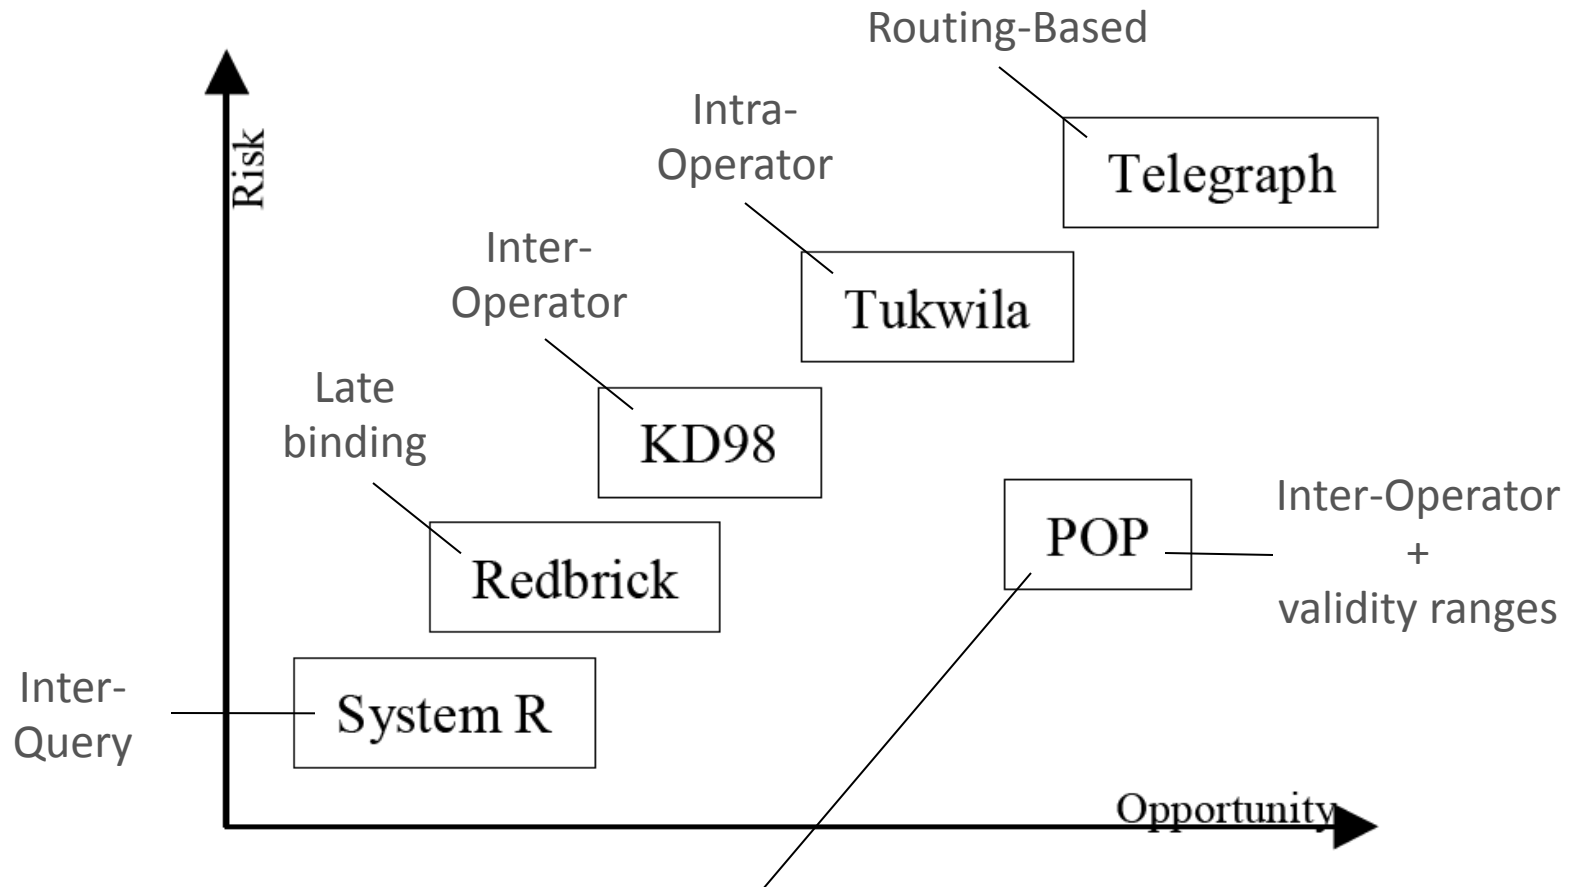
S.b=T.b AND

T.c=U.c





Trade-off between Risk and Re-Optimization Opportunity




[Volker Markl, Vijayshankar Raman, David E. Simmen, Guy M. Lohman, Hamid Pirahesh: Robust Query Processing through Progressive Optimization. SIGMOD 2004:659-670]




Both use validity ranges to detect suboptimalities but initiated differently

Reactive Reoptimization

1) Compute Validity Range (BLACK BOX)



2) Place Checkpoint operators




3) Trigger Re-Optimization if validity range is violated at checkpoint operator

Reactive


Proactive Reoptimization

1) Bounding box around all estimates



Proactive

2) Use bounding boxes to compute a switchable plan for certain statistic ranges (lo, est, hi)



3) Switch Operator chooses path of the plan

R



Plan-Based Adaptation (DBMS)

- Reoptimization scope: current plan only (future queries will not benefit)

Continuous-Query-Based Adaptation (DSMS)

- Reoptimization scope: continuous query (future incoming tuples benefit)
- Temporal and state management aspects

Routing-Based Adaptation (DBMS, DSMS)

- Query processing and reoptimization as routing of tuples along different paths

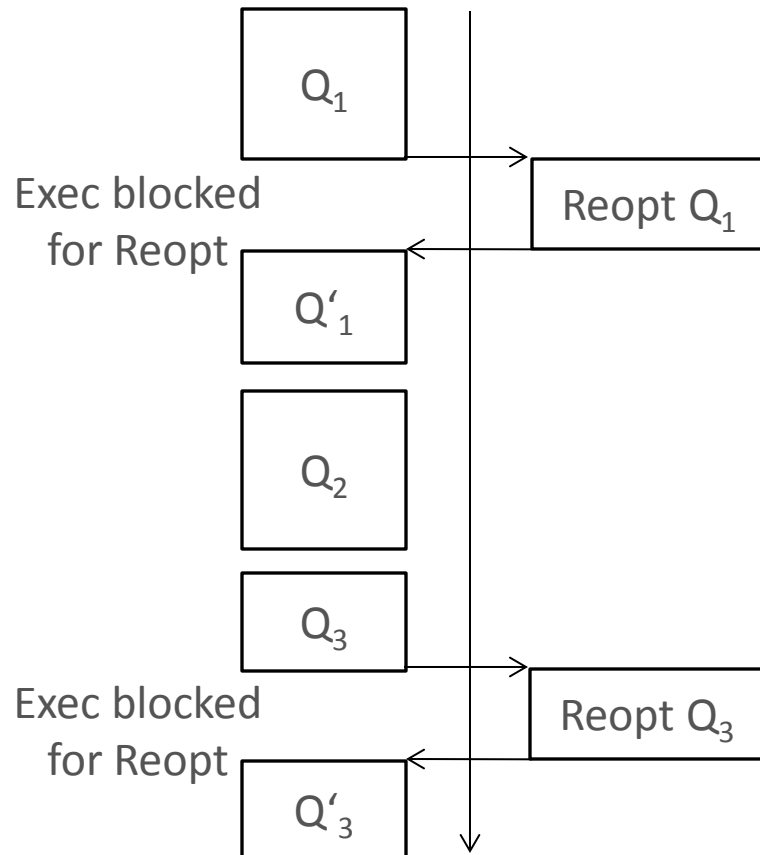
Integration-Flow Optimization (IS)

- Many independent instances (deployed once, executed many times)
- Periodical reoptimization (independent of executing a single instance)
- On-demand reoptimization
 - Model optimality conditions of a plan and continuously check if optimality conditions are violated




Synchronous versus Asynchronous Re-Optimization

Synchronous Reopt
(e.g., mid-query optimization)



Asynchronous Reopt
(e.g., reoptimization in DSMS/IS)



Typical granularity: inter-query



Lesson Learned

- Temporal Classification
(Inter-Query, Late Binding, Inter-Operator, Intra-Operator, Routing -Based)
- Reactive versus Proactive Re-Optimization
- Rewriting-Oriented Classification
- Synchronous versus Asynchronous Re-Optimization



Advantages

- Reoptimization Opportunity: Adaptation to misestimated or unknown properties
➔ Huge execution time improvement

Disadvantages

- Risk of Overhead: Overhead for evaluating optimality, re-optimization and reuse of intermediate results
➔ Overhead might not be amortized by execution time improvement



◆ CLASSIFICATIONS OF ADAPTIVE QUERY PROCESSING

- ❖ Inter-Query Adaptation (static, late binding)
- ❖ Inter-Operator Adaptation
- ❖ Intra-Operator Adaptation
- ❖ Tuple Routing
- ❖ Rewriting-Based Classification
- ❖ Reactive vs. Proactive Reoptimization
- ❖ Synchronous vs. Asynchronous Reoptimization

◆ ADAPTIVE QUERY PROCESSING (IN ACTION)

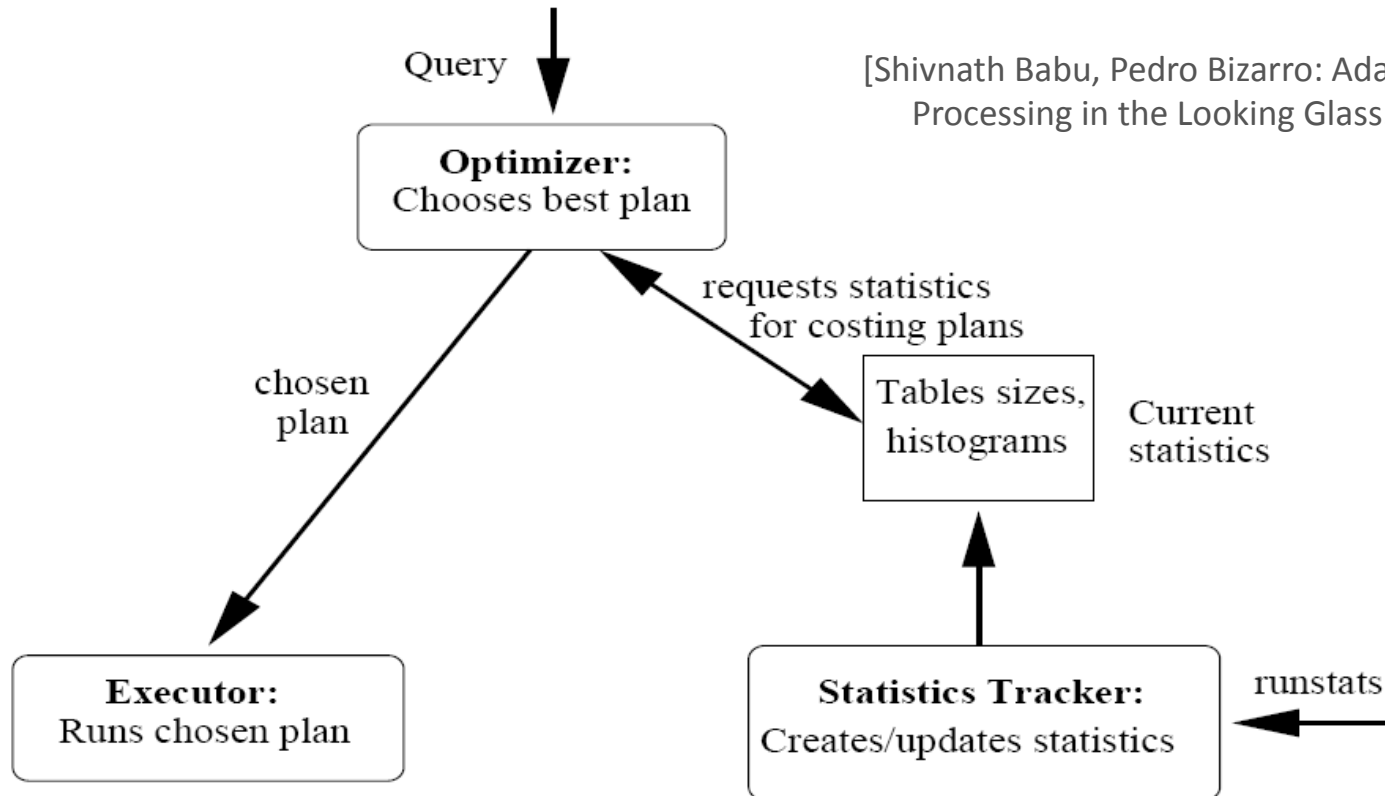
- ❖ Plan-Based Adaptation
- ❖ Continuous-Query-Based Adaptation
- ❖ Routing-Based Adaptation



Adaptive Query Processing (in Action)



Starting Point: Traditional DBMS (inter-query)



Interesting remark:

[Surajit Chaudhuri: Query optimizers: time to rethink the contract? SIGMOD 2009]

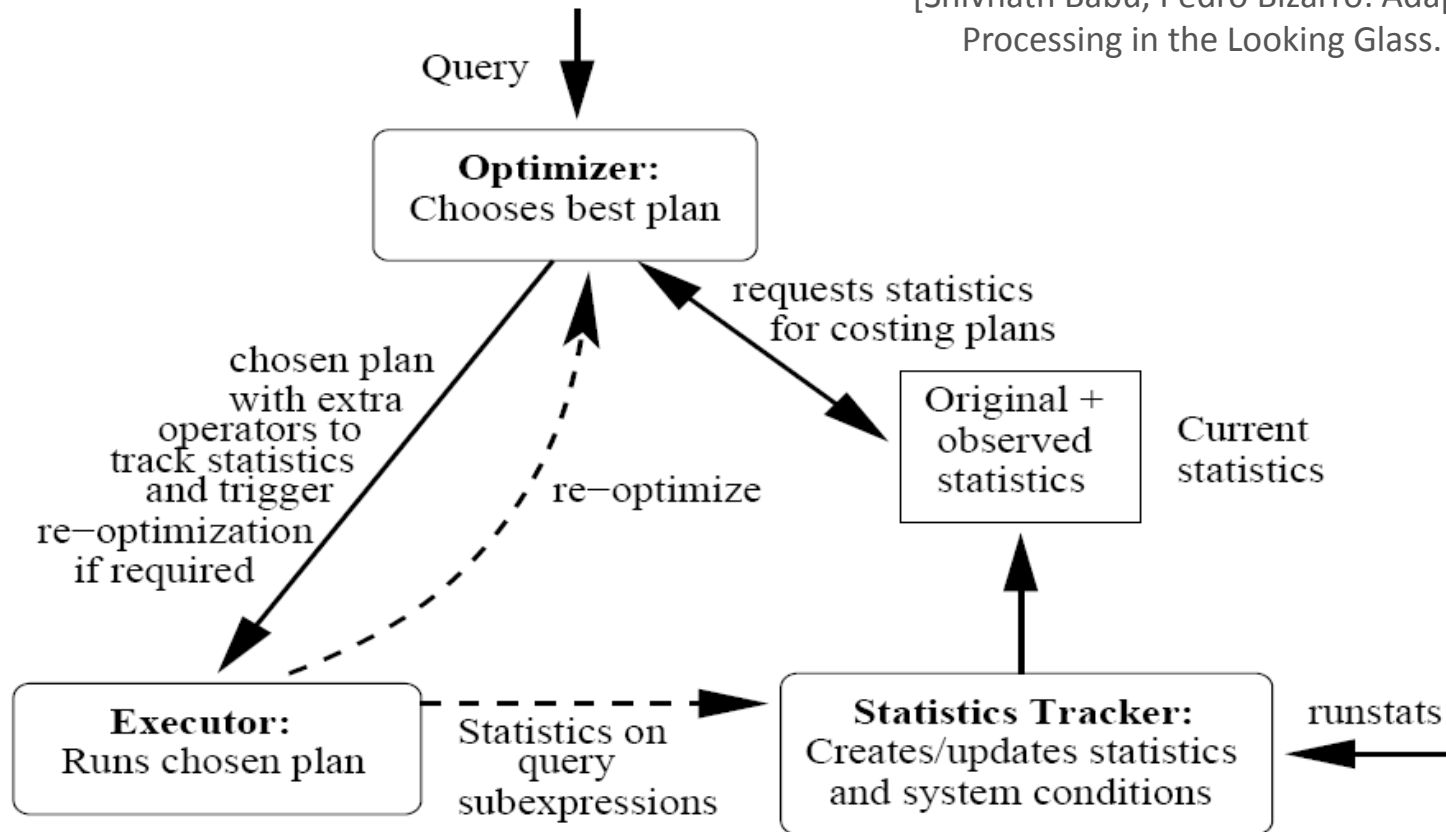


Plan-Based Adaptation



Reoptimize the current plan during execution

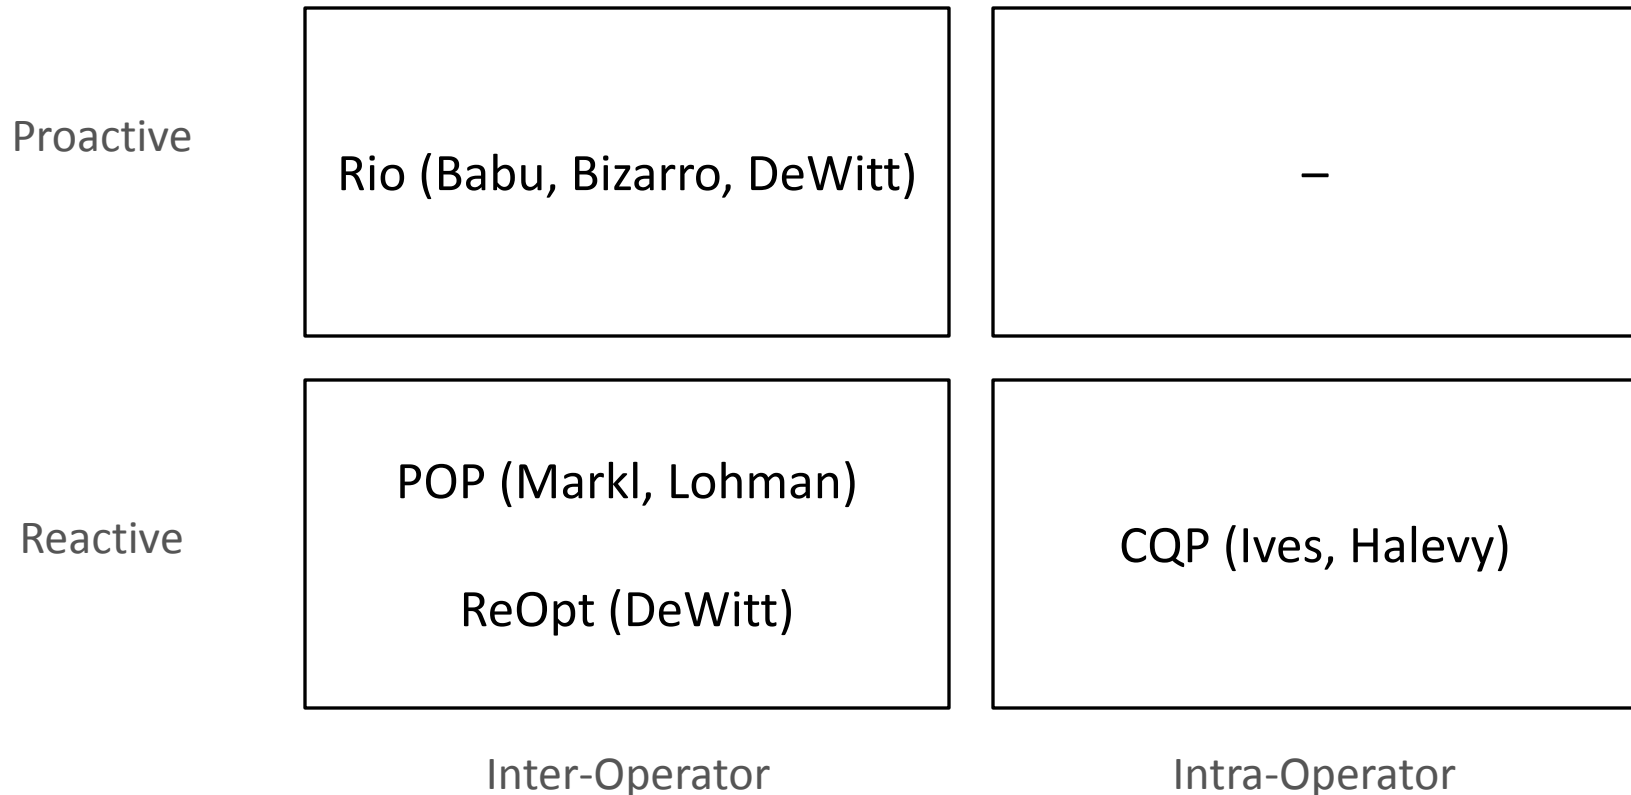
[Shivnath Babu, Pedro Bizarro: Adaptive Query Processing in the Looking Glass. CIDR 2005]





Important plan-based approaches

- All **inter-operator** approaches are **synchronous**, while CQP uses **asynchronous** reoptimization

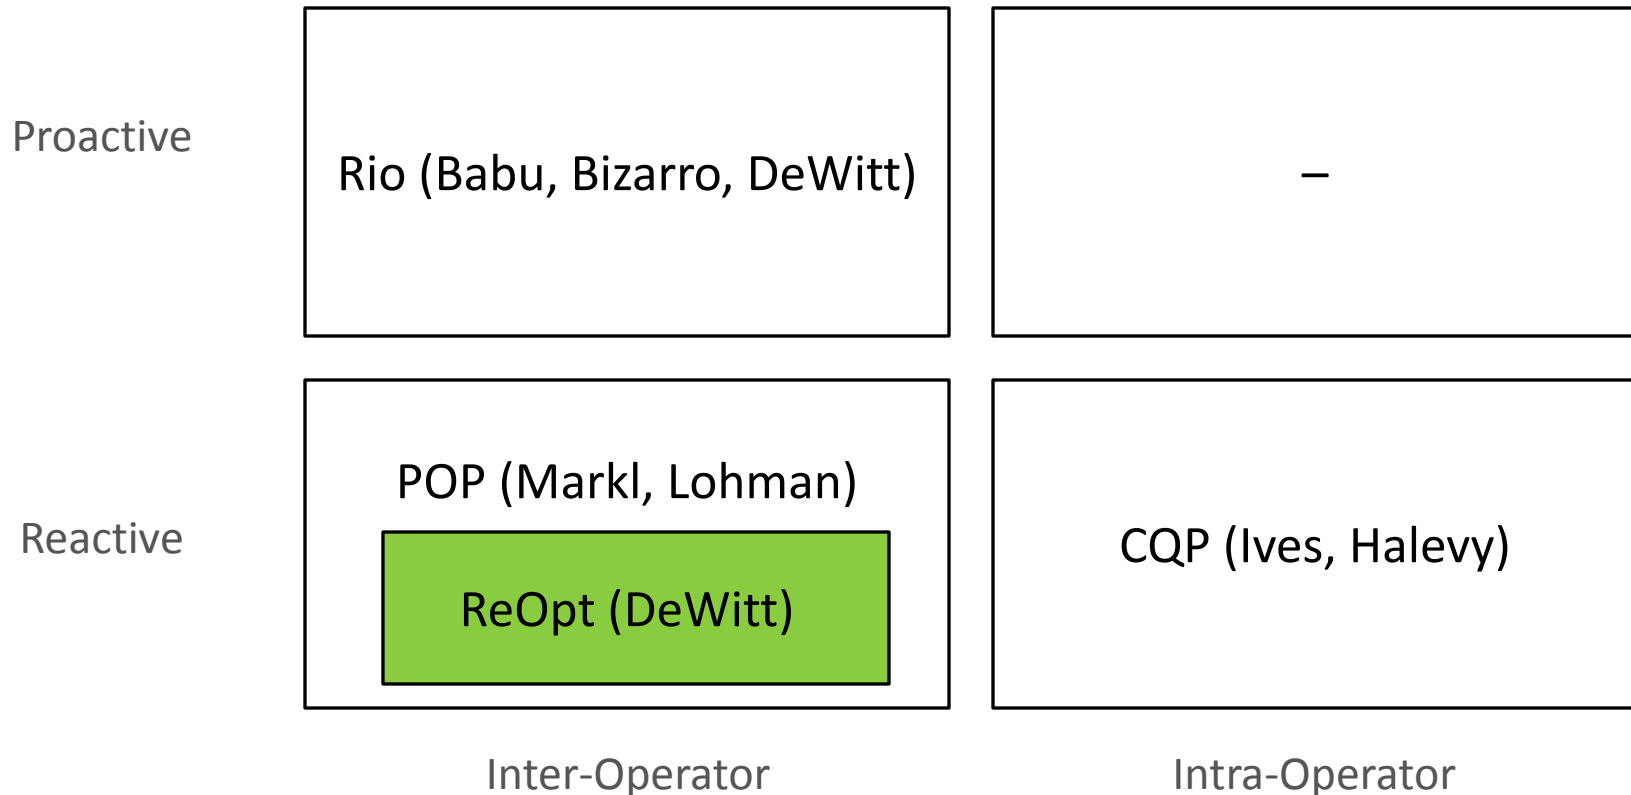




[Navin Kabra, David J. DeWitt: Efficient Mid-Query Re-Optimization of Sub-Optimal Query Execution Plan, SIGMOD 1998]

Important plan-based approaches

- All **inter-operator** approaches are **synchronous**, while CQP uses **asynchronous** reoptimization





Solution Overview

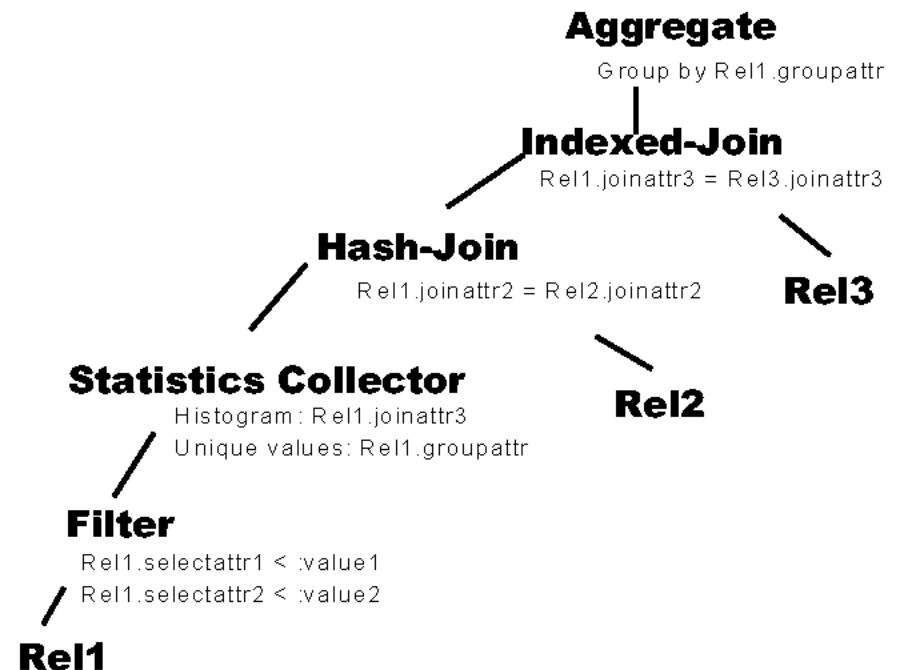
- monitor how the query is doing at key points, and consider dynamically re-optimizing those portions of the query which have not yet been started
- Using System-R optimizer

Core Aspects

- Annotated Query Execution Plans
- Runtime Collection of Statistics
- Dynamic Resource Re-allocation
- Query Plan Modification

Annotated Query Plans

- Statistic monitoring
 - Sizes and cardinalities, selectivities of predicates
 - Estimates of number of groups to be aggregated
- Add into tree statistic collectors
 - Must be collectable in a single pass
 - Will only help with portions of query “beyond” the current pipeline





Low-Overhead Statistics

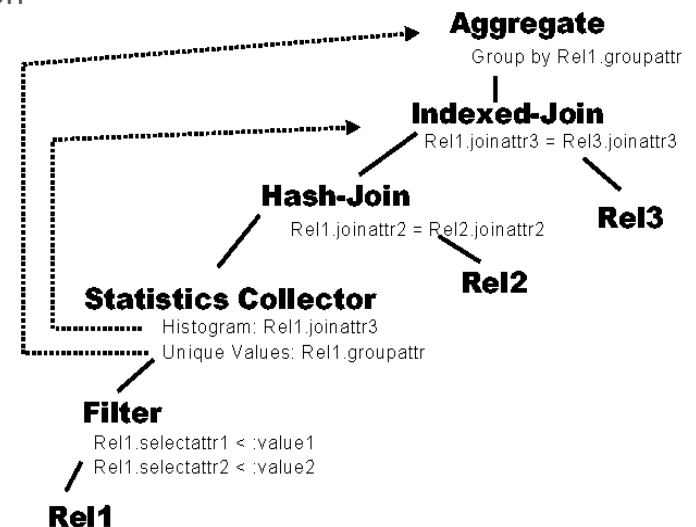
- Want to find “most effective” statistics
 - Don’t want to gather statistics for “simple” queries
 - Want to limit effect of algorithm to maximum overhead ratio, μ
 - Factors: Probability of inaccuracy, Fraction of query affected

Inaccuracy Potentials

- The following heuristics are used:
 - Inaccuracy potential = low, medium, high
 - Lower if we have more information on table value distribution
 - 1+max of inputs for multiple-input selection
 - Always high for user-defined methods
 - Always high for non-equi joins
 - For most other operators, same as worst of inputs

More Heuristics

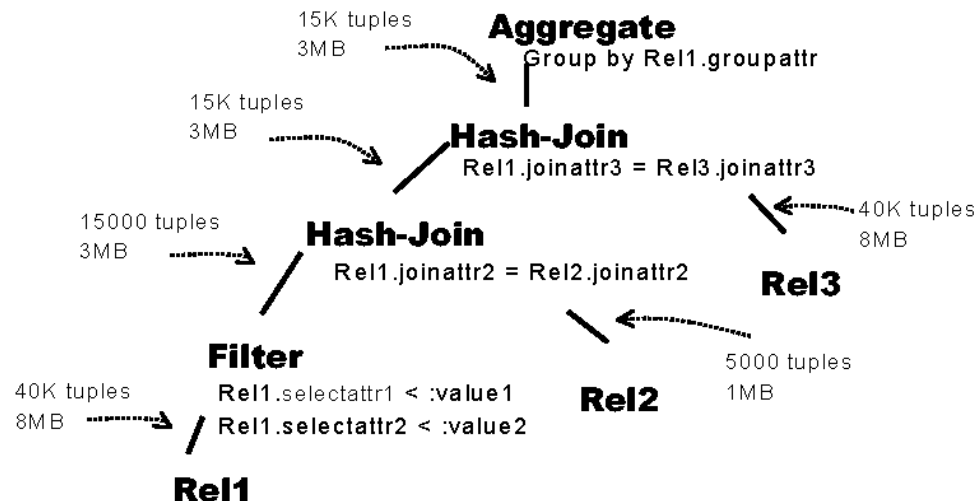
- Check fraction of query affected
 - Check how many other operators use the same statistic
- The winner:
 - Higher inaccuracy potentials first
 - Then, if a tie, the one affecting the larger portion of the plan





*Based on improved estimates,
we can modify the memory allocated
to each operation*

- Results: less I/O, better performance
- Only for operations that have not yet begun executing

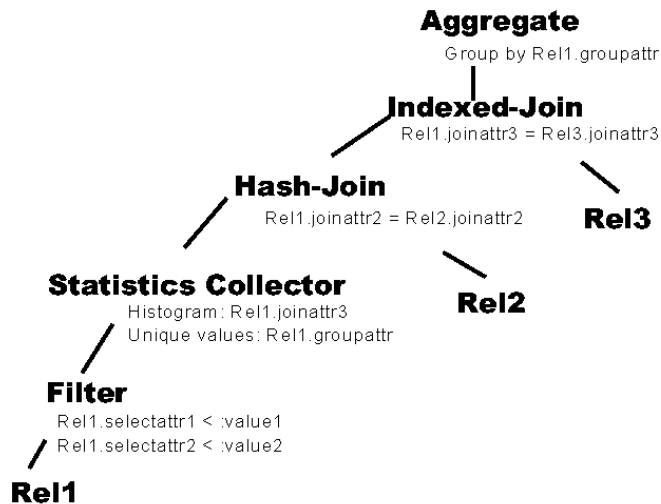




- 1) Only re-optimize part not begun
- 2) Suspend query, save intermediate in temp file
- 3) Create new plan for remainder, treating temp as an input

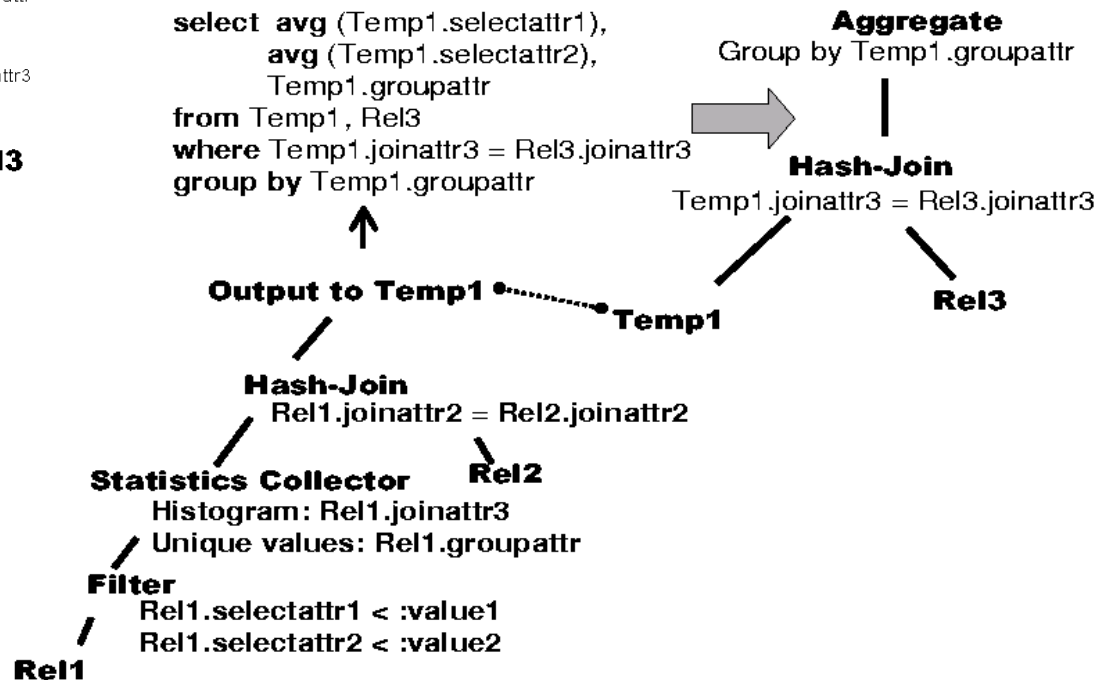
Discussion:
ad-hoc thresholds are a bad idea

- E.g. even a 100x error on very small relation may not make a difference in optimal plan
- Many unnecessary re-optimization steps



```

select avg (Temp1.selectattr1),
       avg (Temp1.selectattr2),
       Temp1.groupattr
from Temp1, Rel3
where Temp1.joinattr3 = Rel3.joinattr3
group by Temp1.groupattr
    
```

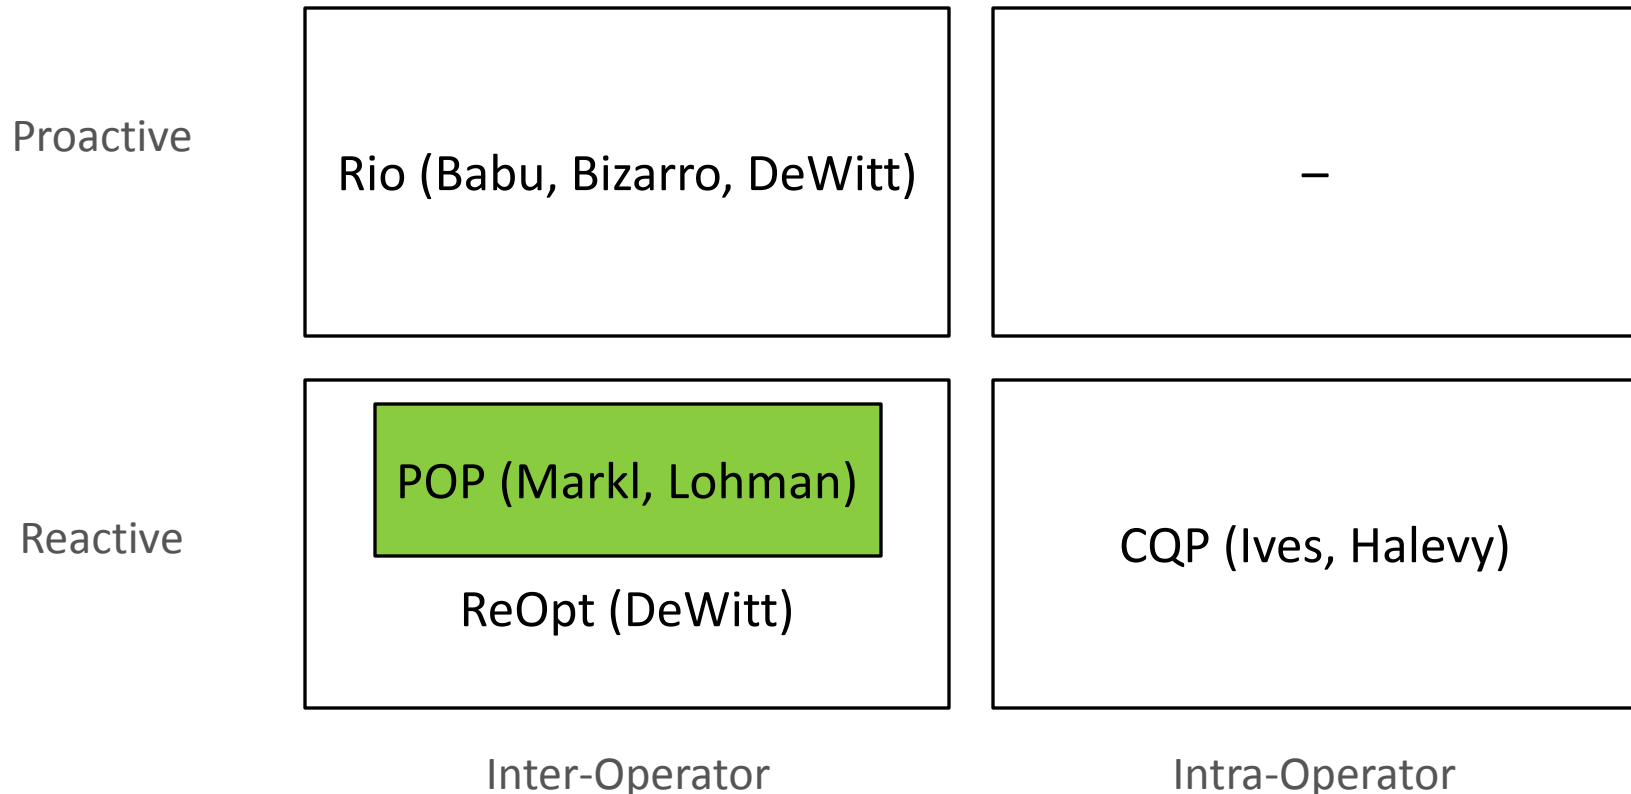




[Volker Markl, Vijayshankar Raman, David E. Simmen, Guy M. Lohman, Hamid Pirahesh: Robust Query Processing through Progressive Optimization, SIGMOD 2004]

Important plan-based approaches

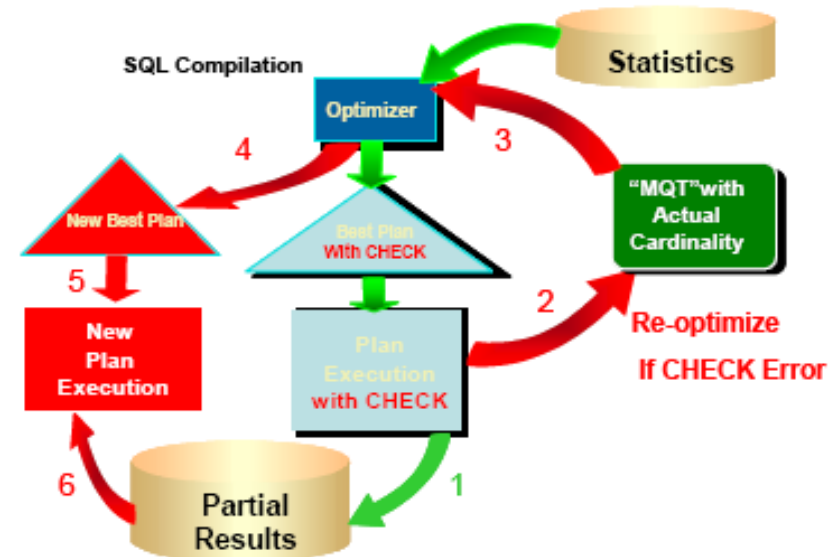
- All **inter-operator** approaches are **synchronous**, while CQP uses **asynchronous** reoptimization



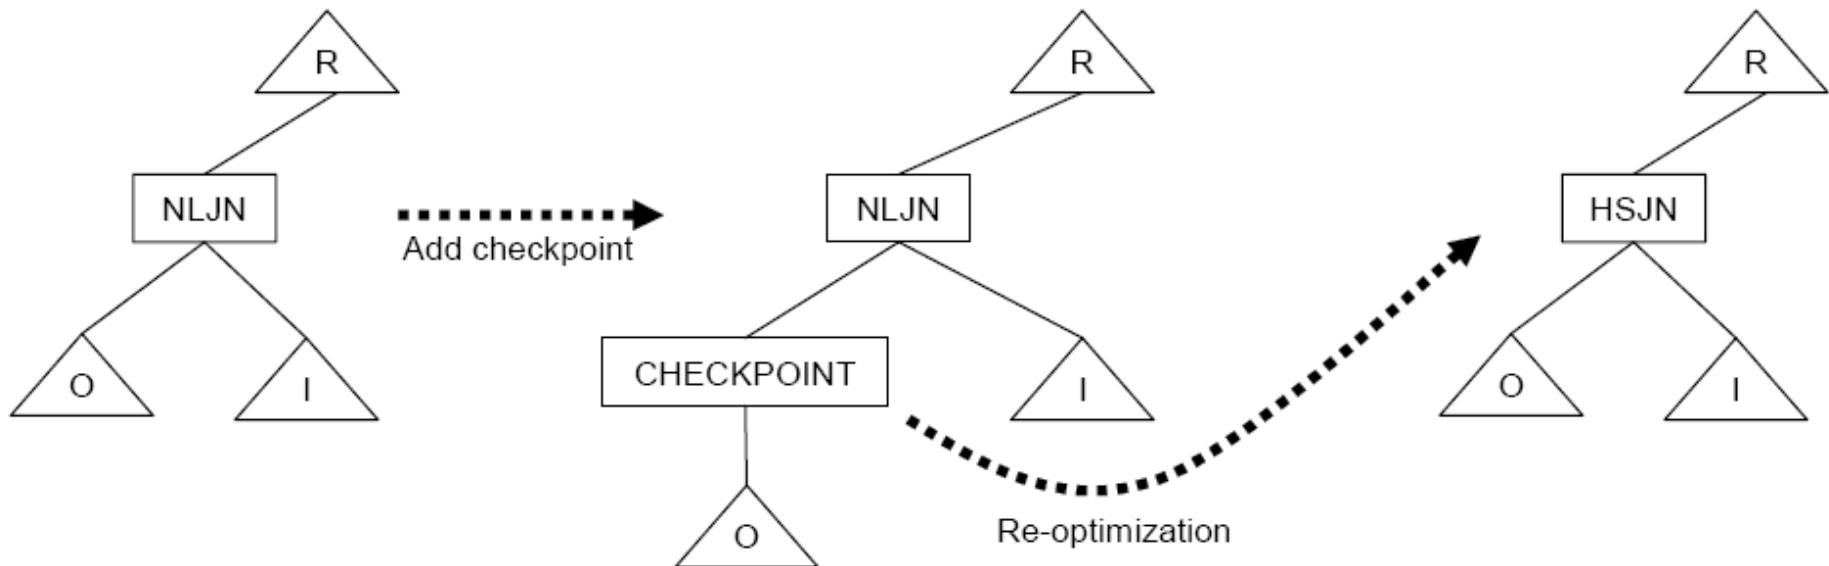


Solution Overview

- Lazily trigger reoptimization during execution if cardinality counts indicate current plan is suboptimal
- introduces checkpoint (CHECK) operator to compare actual vs estimated cardinality
- precompute optimal cardinality ranges



Risk vs. Opportunity





CHECK operator to find if a plan is suboptimal

- At optimization time, find out cardinality range (at CHECK location) for which plan is optimal
- At run time, ensure cardinality within $[l, u]$
- If violated, stop plan execution and reoptimize

Location of CHECKs

Re-optimize

- taking observed cardinality into account, and
- exploiting intermediate results where beneficial
- Heuristic: limit number of reoptimizations (default: 3)

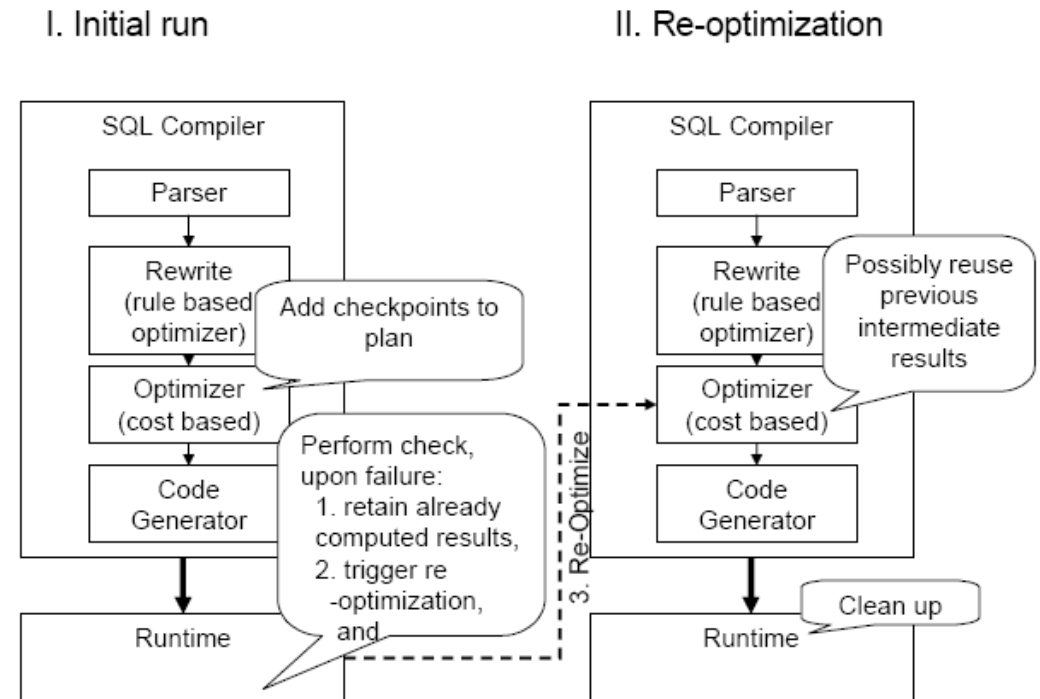


Figure 1: Progressive Optimization architecture



Consider a plan edge e that flows rows into operator o ,

- let P be the subplan rooted at o .
- The **validity range** for e is an upper and lower bound on the number of rows flowing through e , such that if the range is violated at runtime, we can guarantee P is suboptimal

Finding Optimality Ranges

- Plan P_{opt} with root operator o_{opt} is being compared with another plan P_{alt} different only in the root operator o_{alt} .
- Need to solve
 - $\text{cost}(P_{\text{alt}}, c) - \text{cost}(P_{\text{opt}}, c) = 0$
 - where c is the cardinality on edge e
- Cost functions can be complex/non-linear/non-continuous
- Using Newton-Raphson Iteration

Conservative Detection of Suboptimalities




Figure 4: Computing the Upper Bound of a Validity Range

Discussion:

black-box approach for the sub query plan

➔ No optimality conditions

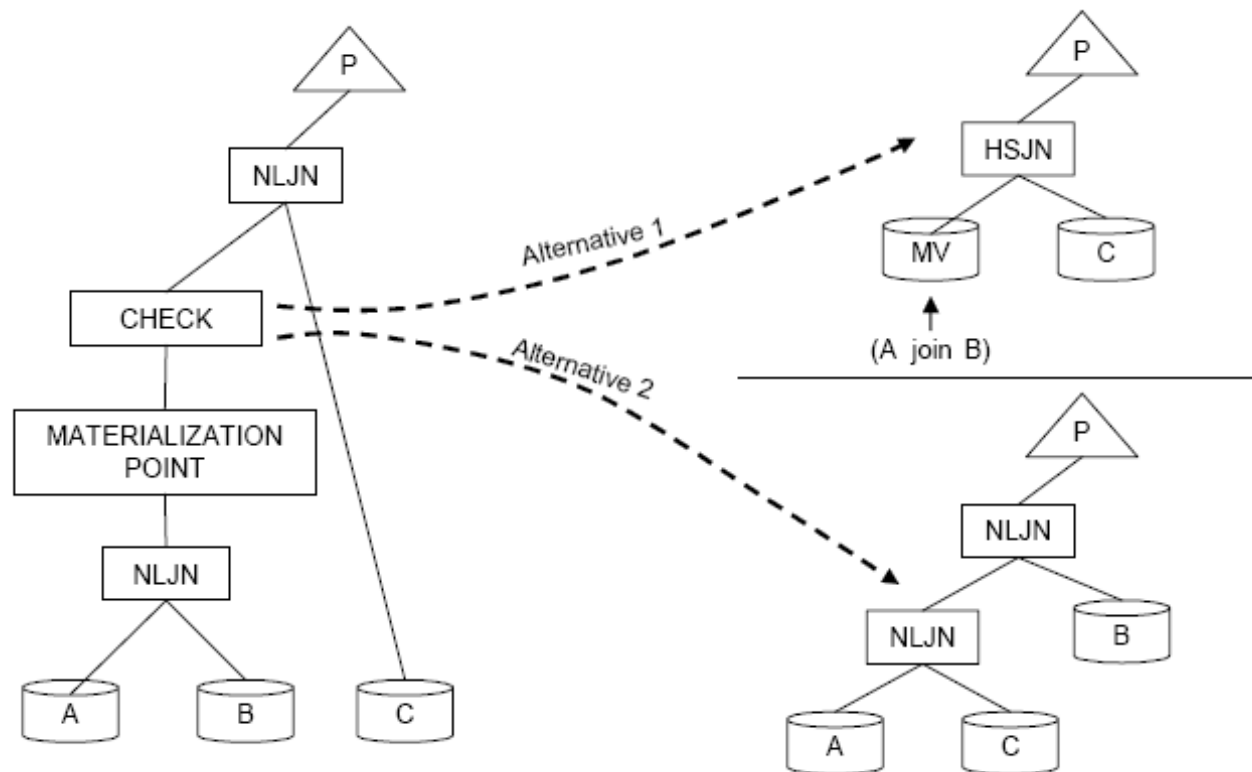
➔ No directed reoptimization

➔ Validity ranges not combined for multiple sub graphs?



Optional Use of Materialized Views

- But all intermediate results are stored as temporary MVs
- Optimizer decides to use them or not (trashing or reusing as optimization problem)

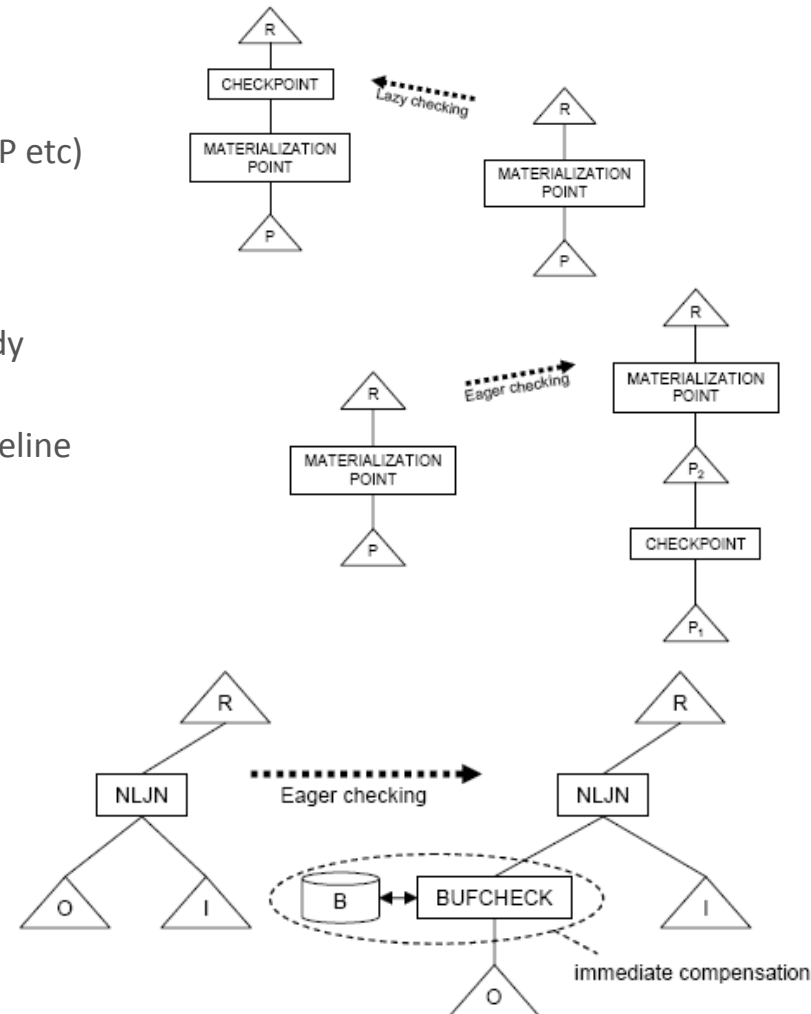




Variants applicable in different cases, trade off risk for opportunity

Variants

- Lazy checking
 - Adding CHECKs above a materialization point (SORT, TEMP etc)
- Lazy checking with eager materialization
 - Can insert materialization point if it does not exists already
- Eager checking without compensation
 - CHECK is pushed down the materialization point, into pipeline
- Eager checking with buffering
 - Delayed pipelining
- Eager checking with deferred compensation
 - Anti-join all rows returned to the user with the new result stream

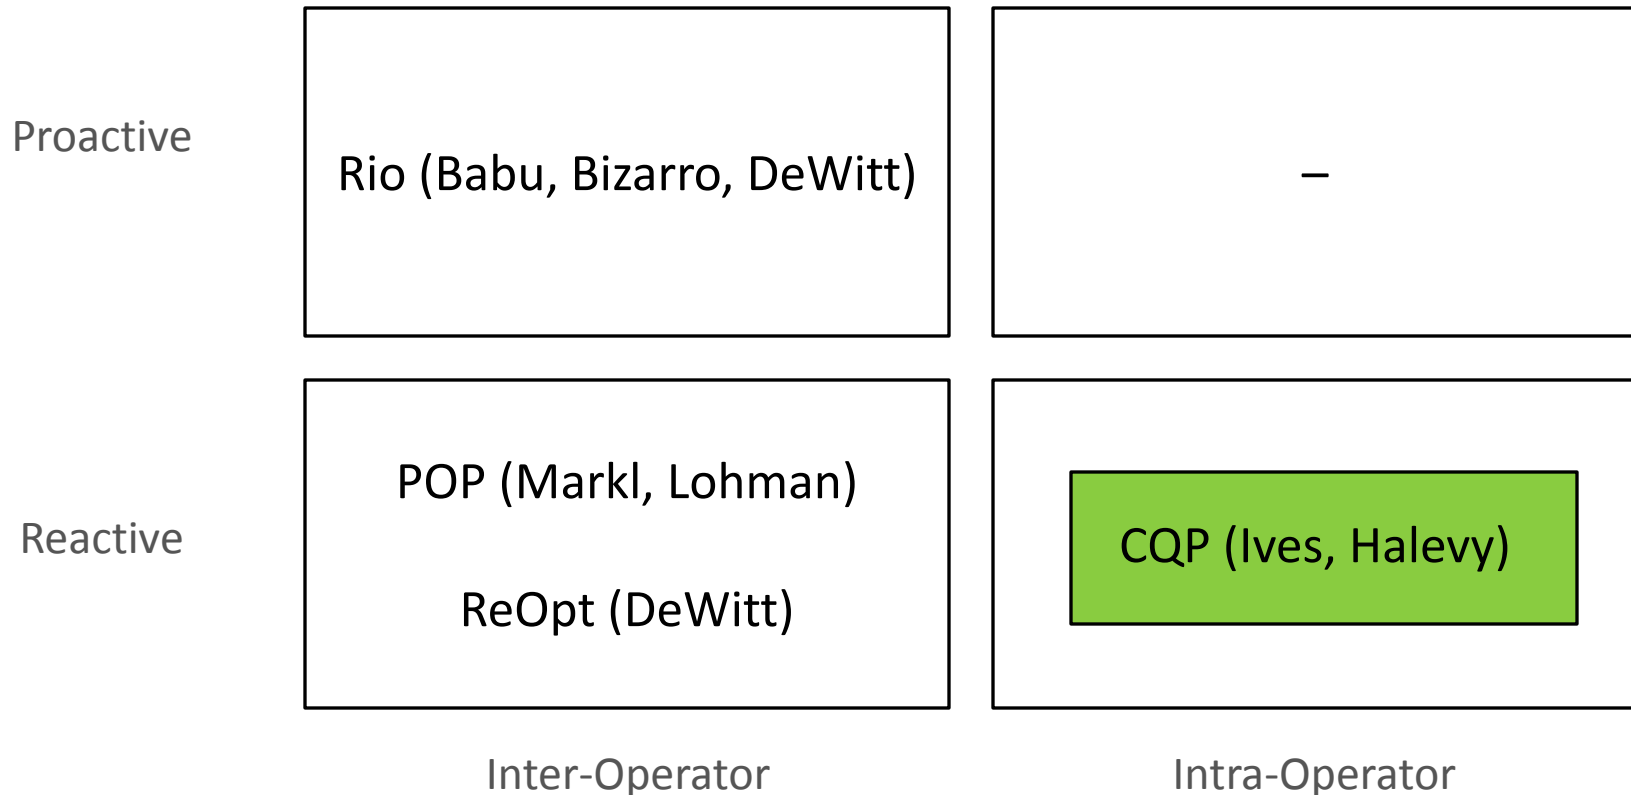




[Zachary G. Ives, Alon Y. Halevy, Daniel S. Weld:
Adapting to Source Properties in Processing Data
Integration Queries. SIGMOD 2004]

Important plan-based approaches

- All **inter-operator** approaches are **synchronous**, while CQP uses **asynchronous** reoptimization





Drawbacks of Inter-Operator Re-Optimization


- Possibly trashing of intermediate results
- Use of materialization points might be too coarse-grained

Core Idea (within the data integration platform Tukwila)

- Data partitioning – use new plans for new data only
- Combine resulting data partitions in so-called stitch-up phases

History:

Convergent Query Processing
renamed to
Corrective Query Processing



Flights
F(fid,from,to,when)

				F ⁰
				F ¹

Travelers
T(ssn,flight)

		T ⁰
		T ¹

Number of children per traveler
C(parent,num)

		C ⁰
		C ¹

> Pipelined Hash Join



Pipelined Hash Join

- symmetric
- non-blocking
- memory-consuming
- foreach arriving tuple (P1 | P2)
 - 1) Probing
 - 2) Building

fid	from	to	when	ssn
1007	Berlin	Chicago	3456712	70001
1007	Berlin	Chicago	3456712	70002
1001	Dresden	LA	1234567	70003
1001	Dresden	LA	1234567	70004

Hashtable 1	
	1001 1007

Hashtable 2	
	1007 1007 1001 1001

fid=flight

Flight

fid	from	to	when
1001	Dresden	LA	1234567
1007	Berlin	Chicago	3456712

P1

Travelers

ssn	flight
70001	1007
70002	1007
70003	1001
70004	1001

P2



Algebraic key property

- Distribute relational UNION through PSJ operations

- For joins; in general:
$$R_1 \bowtie \dots \bowtie R_m = \bigcup_{1 \leq c_1 \leq n, \dots, 1 \leq c_m \leq n} (R_1^{c_1} \bowtie \dots \bowtie R_m^{c_m})$$
- Example:

If $R_1 = R_1^1 \cup R_1^2$, (horizontal partitioning)
 $R_2 = R_2^1 \cup R_2^2$

then: $R_1 \bowtie R_2$

$$\equiv (R_1^1 \cup R_1^2) \bowtie (R_2^1 \cup R_2^2)$$

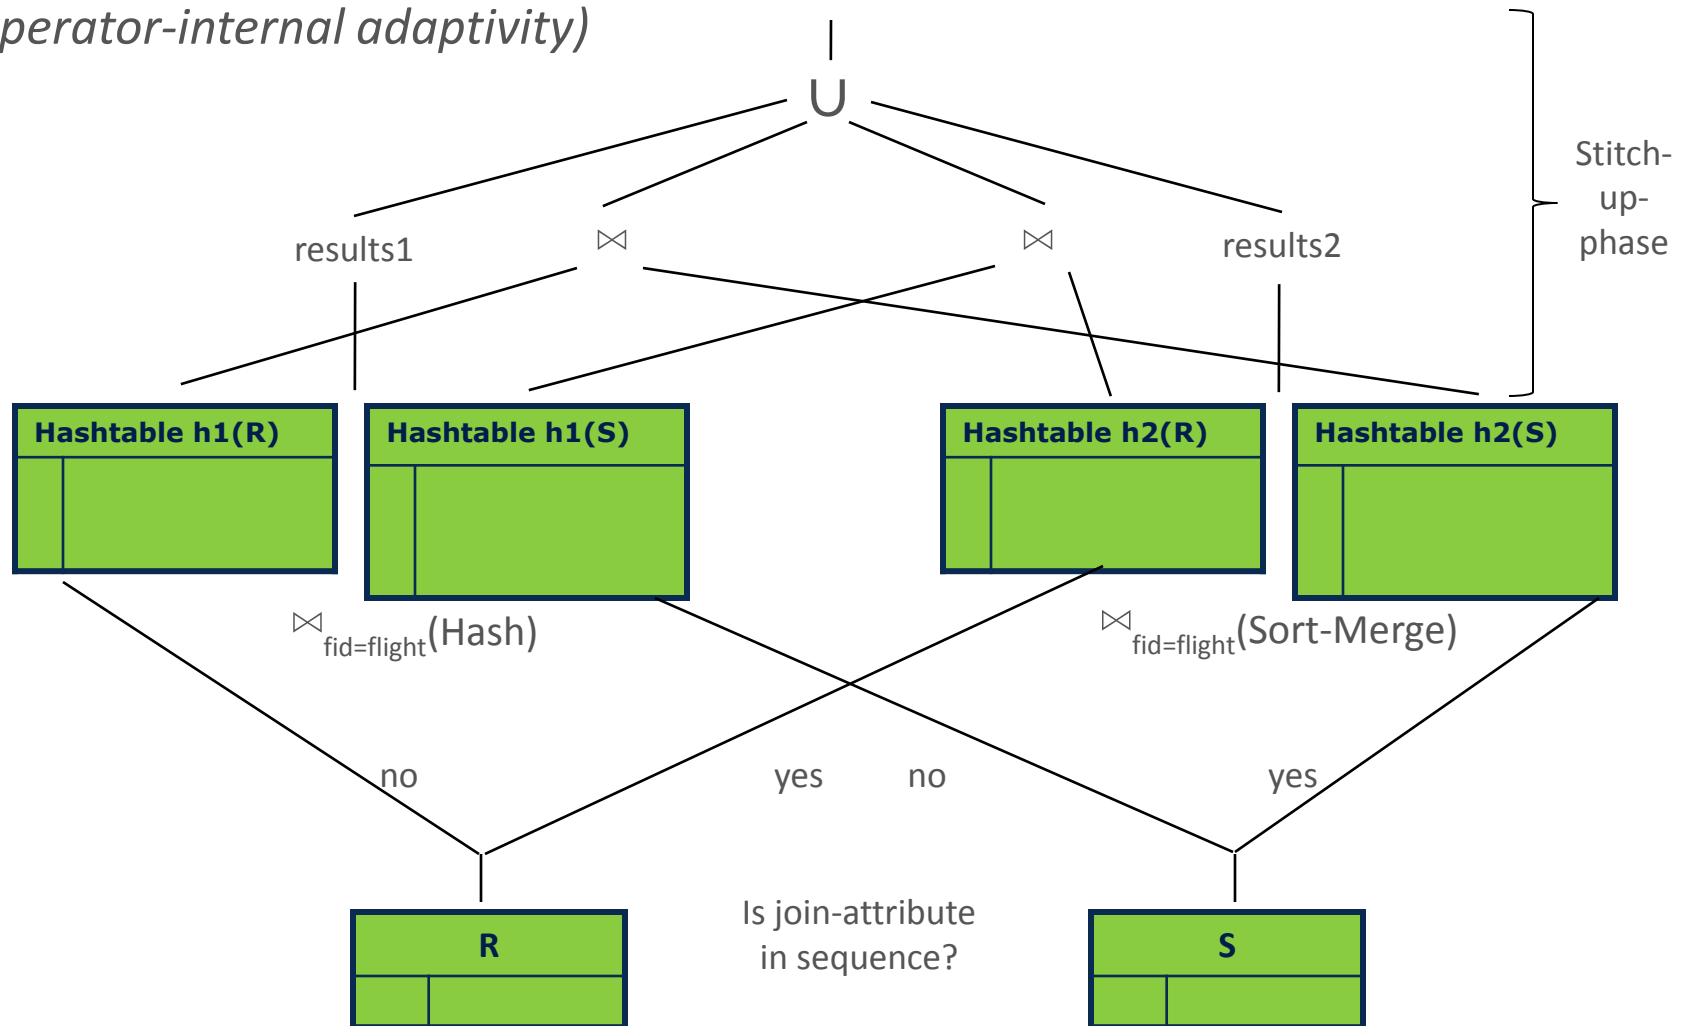
$$\equiv (R_1^1 \bowtie R_2^1) \cup (R_1^2 \bowtie R_2^2) \cup (R_1^1 \bowtie R_2^2) \cup (R_1^2 \bowtie R_2^1)$$



Union over all join combinations of partitions

(state in terms of results and hash maps)

*Complementary Join Pair
(operator-internal adaptivity)*





Eager Group By




- Cost-based rewriting decision

Adjustable-window pre-aggregation

- Pipelined pre-aggregation
- Window required for pipeline character

Control Strategy

- if current window is effective, increase next window size
- if current window is not effective, reduce next window size (until 1)
- used before
 - Join
 - Final-Aggregation

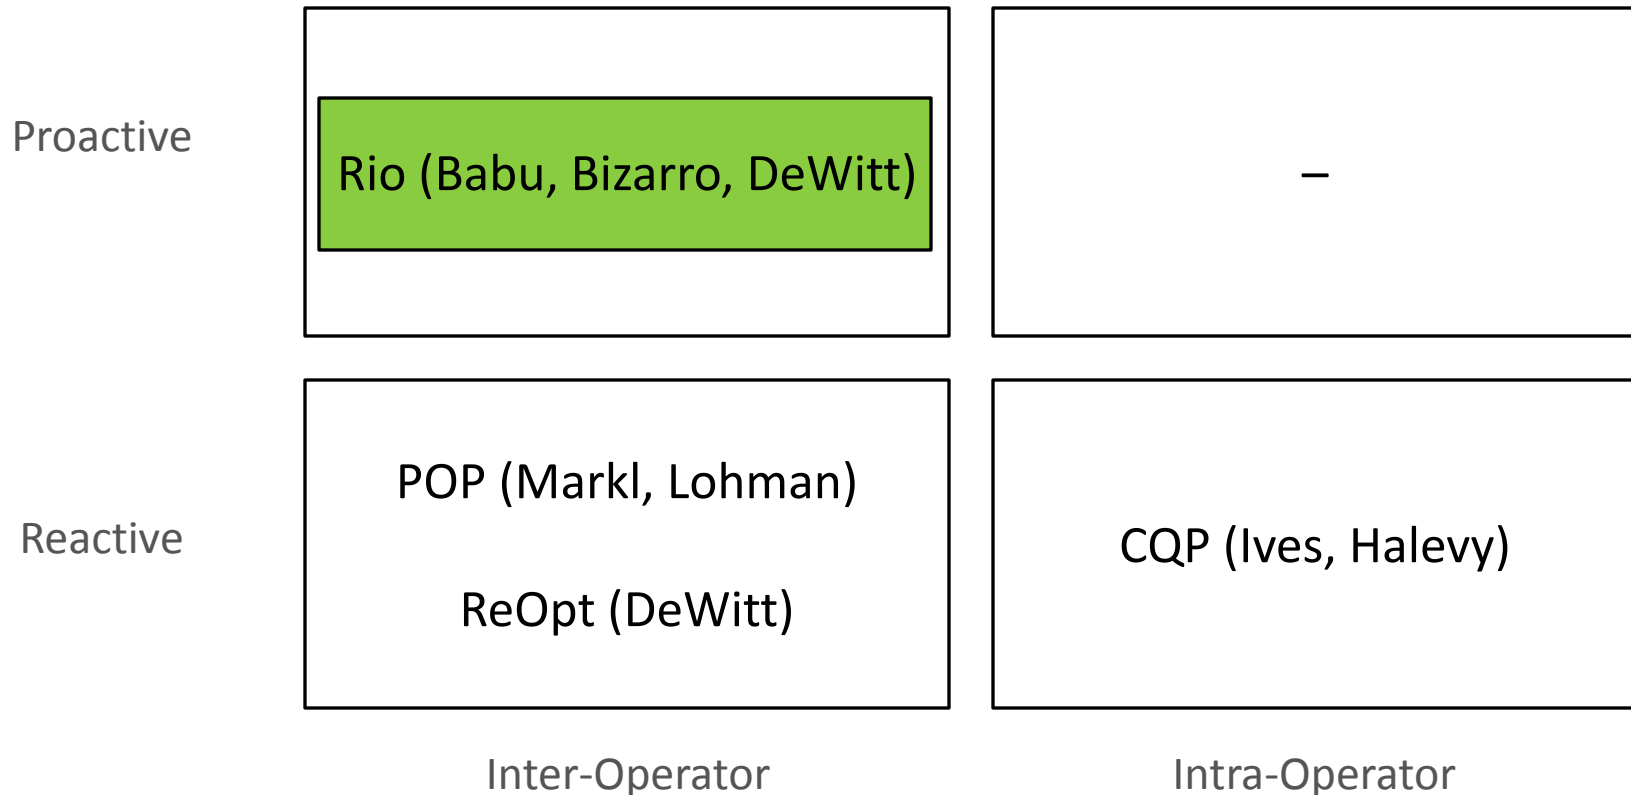




[Shivnath Babu, Pedro Bizarro, David J. DeWitt:
Proactive Re-optimization. SIGMOD 2005]

Important plan-based approaches

- All **inter-operator** approaches are **synchronous**, while CQP uses **asynchronous** reoptimization





Current Re-Optimizers (reactive)

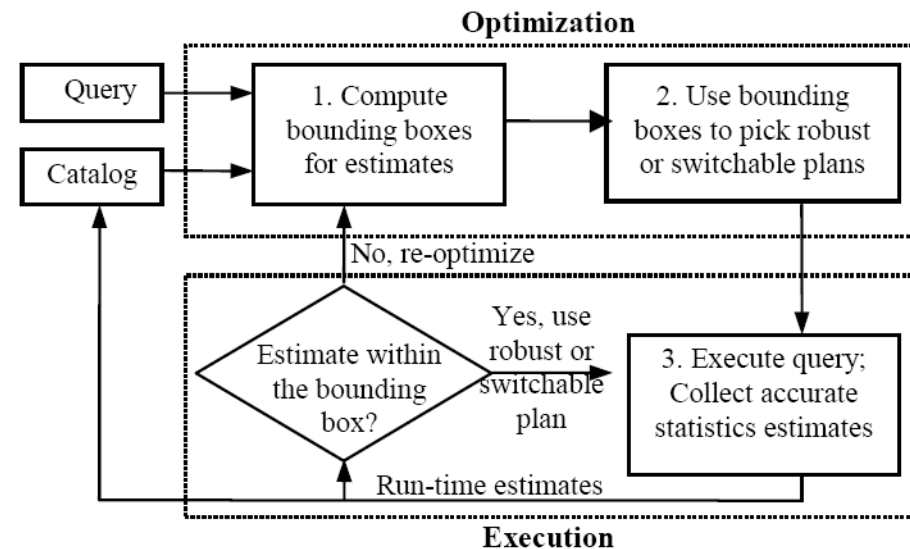
- use traditional optimizers to pick a plan and
- if violations are detected, they stop execution, and re-optimize

Shortcomings of the use of traditional query optimizers (TRADs)

- **Pr1:** pick plans depending on uncertain statistics, making *re-optimization very likely*
- **Pr2:** if TRAD chooses a new plan, *work until estimation error is found will be lost*
- **Pr3:** limited statistics collection ability *leading to new mistakes*

Construction of proactive re-optimizer (Rio)

- *bounding boxes*
(intervals instead of single-point),
- generate *robust* and *switchable* plans
- minimizing re-optimization
- Reuse intermediate in case of plan change
- using *Random Sample processing*





Example 1

- DB-cache=200MB, $|R|=500\text{MB}$, $|S|=160\text{MB}$, and $|\sigma(R)|=300\text{MB}$
- optimizer estimates $|\sigma(R)|=150\text{MB}$

Query:


SELECT *

FROM R, S


WHERE R.a=S.a AND

R.b>K1 AND

R.c>K2



- TRAD will choose P1a resulting in making two passes over R and S (P1b needs only one)
- Reactive re-optimizer (VRO) will choose P1a computing a validity range for P1a $100\text{KB} \leq |\sigma(R)| \leq 160\text{MB}$






Example 2

Query:

```
SELECT *
FROM R,S,T
WHERE R.a=S.a AND
S.b=T.b AND
R.c>K1 AND R.d=K2
```

- Assume sizes of the tables are known accurately to be $|R|=200\text{MB}$, $|S|=50\text{MB}$, and $|T|=60\text{MB}$
- Further assume $|\sigma(R)|=80\text{MB}$
- optimizer underestimates as 40KB
- Based on these statistics, the TRAD chooses Plan P3a
- Due to limited information of the actual size of $|\sigma(R)|$, a reactive re-optimizer will trigger re-optimization step-by-step three times moving from P3a to P3d



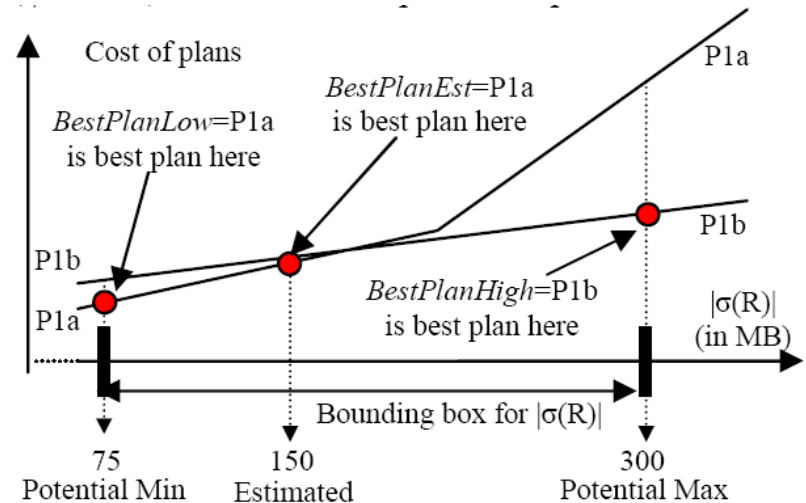




Bounding Boxes and Plans: four cases

- Single optimal plan
- Single robust plan
- Switchable plan (set S of plans p)
- None of the above

Example

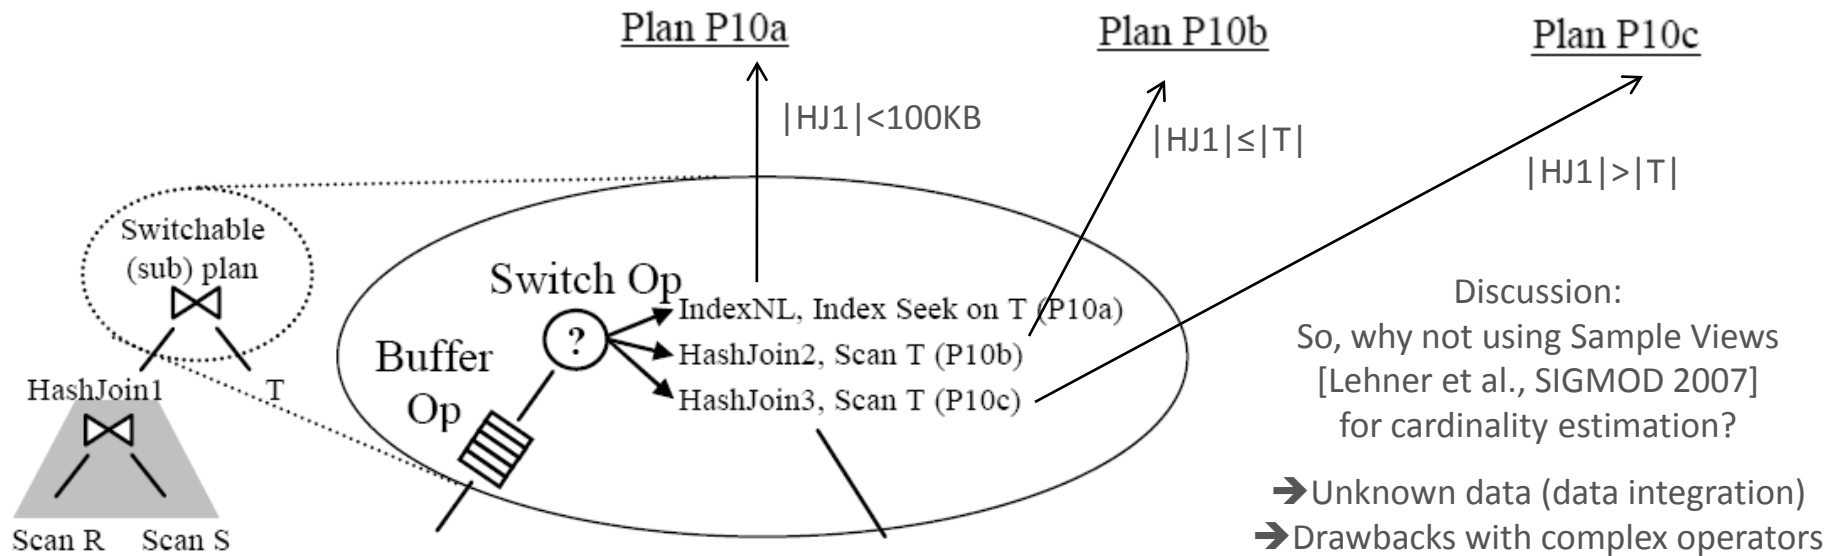

- A) robust plan P1b
- B) switchable plans (always 3 plans)





Switchable Plans

- Reuse intermediate results (no trashing)
- Buffer operator
 - Based on a random sample of its input, it decides which subplan to initiate
 - Random sample punctuations





Lesson Learned

- Reactive, inter-operator reoptimization
- Reactive, inter-operator reoptimization with validity ranges
- Reactive, intra-operator reoptimization
- Proactive, inter-operator reoptimization with bounding boxes



Advantages

- Fast adaptation during runtime of a single query
- Huge opportunities for optimization (unknowns, changing workload)

Disadvantages

- Partly trashing of intermediate results
- Reoptimization does not guarantee to find a better plan
- Overhead for combining/reusing intermediate results

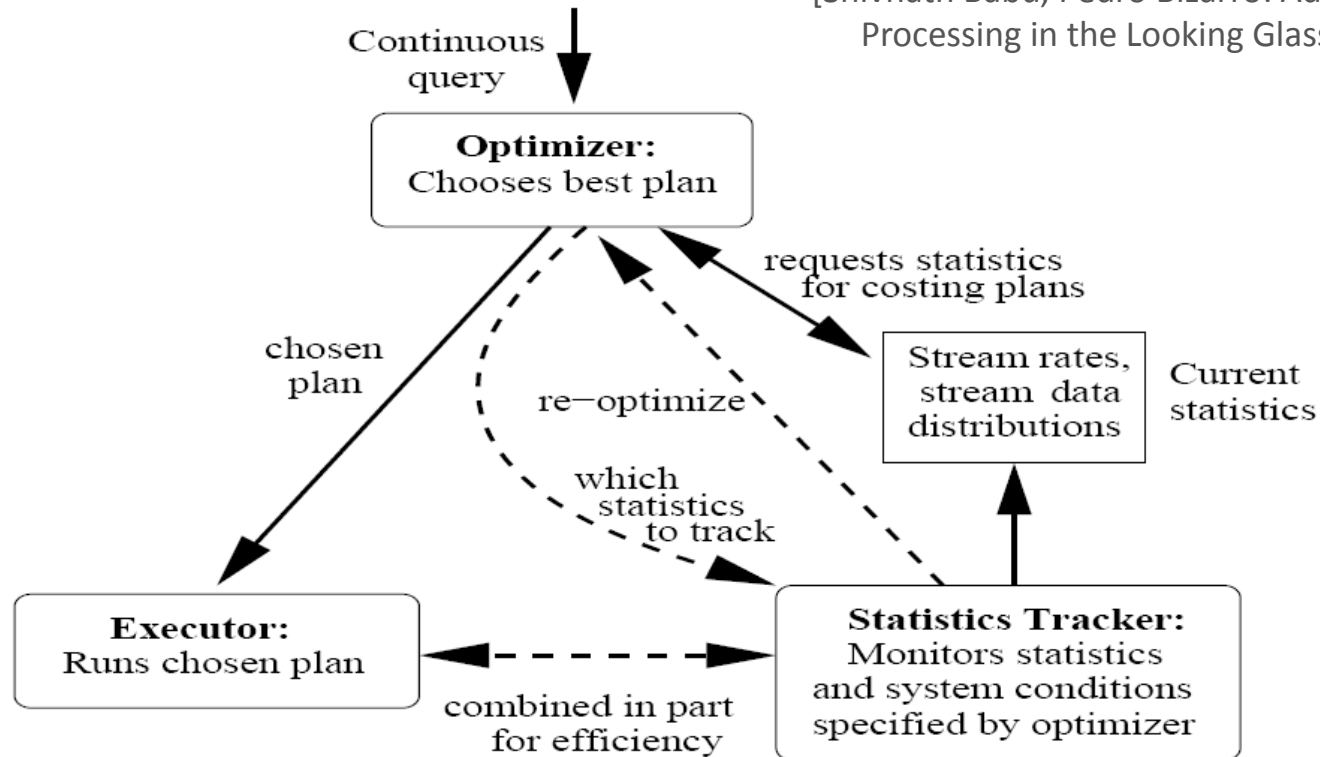


Continuous-Query-Based Adaptation



Reoptimize continuous query during runtime

[Shivnath Babu, Pedro Bizarro: Adaptive Query Processing in the Looking Glass. CIDR 2005]





Fundamental Differences to Mid-Query Re-Optimization

- Larger optimization scope: optimize the continuous query rather than only the current query plan
- Higher optimization potential: all kinds of operators
 - E.g., selection reordering not applicable for plan-based reoptimization but for CQ-based adaptation
- Need for incremental statistics maintenance (no local data)
- Monitor arbitrary changes in stream characteristics and systems conditions during runtime

Example Systems

- CAPE (Worcester)
- NiagaraCQ (Wisconsin-Madison)
- StreaMon (Stanford)

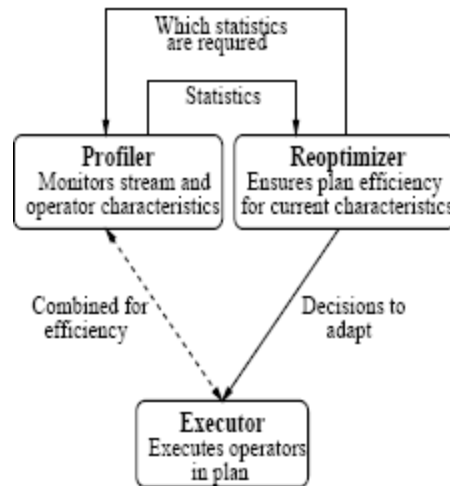
Basic Concepts

- Track only relevant statistics
- Sampling-based techniques
- Combine statistics tracking and query execution whenever possible

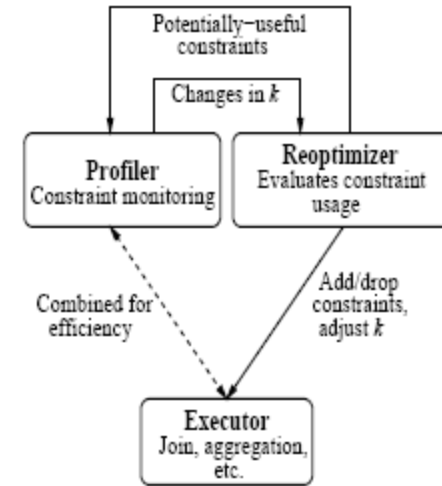
> StreaMon (Stanford)



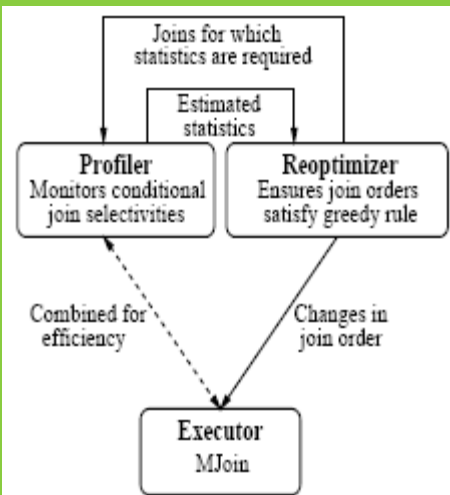
StreaMon
Statistic
Tracking



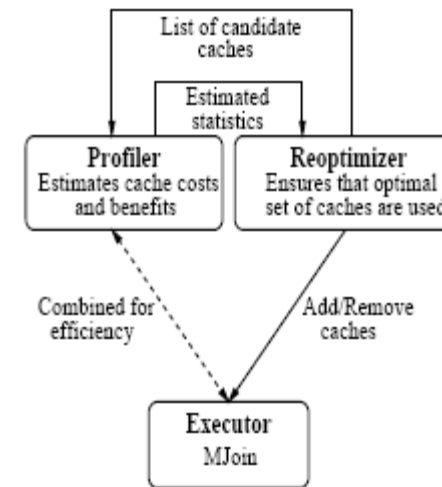
Adaptive
Memory
Minimization



Adaptive
Join Ordering



Adaptive
Caching for
Joins




[Shivnath Babu, Jennifer Widom: StreaMon: An Adaptive Engine for Stream Query Processing. SIGMOD Conference 2004]



[Shivnath Babu, Rajeev Motwani, Kamesh Munagala, Itaru Nishizawa, Jennifer Widom: Adaptive Ordering of Pipelined Stream Filters. SIGMOD 2004]

Pipelined Filters

- Stream of tuples is processed by a set of commutative filters
- E.g., multi-way stream join



(b)

	F_1	F_2	F_3	F_4
	a_1	a_2	a_3	a_4
I	b_1	b_2	b_3	b_4
2	0	0	1	0
4	0	1	1	1
7	1	0	0	0
2	0	0	1	0
5	1	1	0	1
4	0	1	1	1

Processing-time averages

Profile window

(c)

F_3	F_1	F_2	F_4
4	2	3	3
	2	1	1
		0	0
			0

Matrix view V

A-Greedy Algorithm (Adaptive Greedy)

- A-Greedy Profiler: maintains a matrix view over a window of profile tuples
- A-Greedy Re-Optimizer: Checks for violations of greedy invariant and reoptimizes if necessary

(a)

I	b_1	b_2	b_3	b_4
7	1	0	0	0
2	0	0	1	0
5	1	1	0	1
4	0	1	1	1
3	0	0	1	1
6	1	1	0	1

Profile window

(b)

F_3	F_1	F_2	F_4
3	3	3	4
	3	2	2
		0	0
			0

Matrix view V (Violation in first row)

(c)

F_4	F_3	F_1	F_2
4	3	3	3
	1	1	0
		1	0
			0

Matrix view V (After correction)



Lesson Learned

- Adaptation by Dynamic Exploration of Stream Characteristics
- Requirement of Dynamic Plan Migration (Moving/Recomputing State)



[Yali Zhu, Elke A. Rundensteiner, George T. Heineman: Dynamic Plan Migration for Continuous Queries Over Data Streams. SIGMOD 2004]

Advantages

- Huge optimization potential (scope not limited to current query)
- Incremental statistics maintenance required anyway (missing knowledge)

Disadvantages

- Possibly large overhead for full exploration (profiling)
- No evaluation of optimization benefit vs. costs for dynamic plan migration

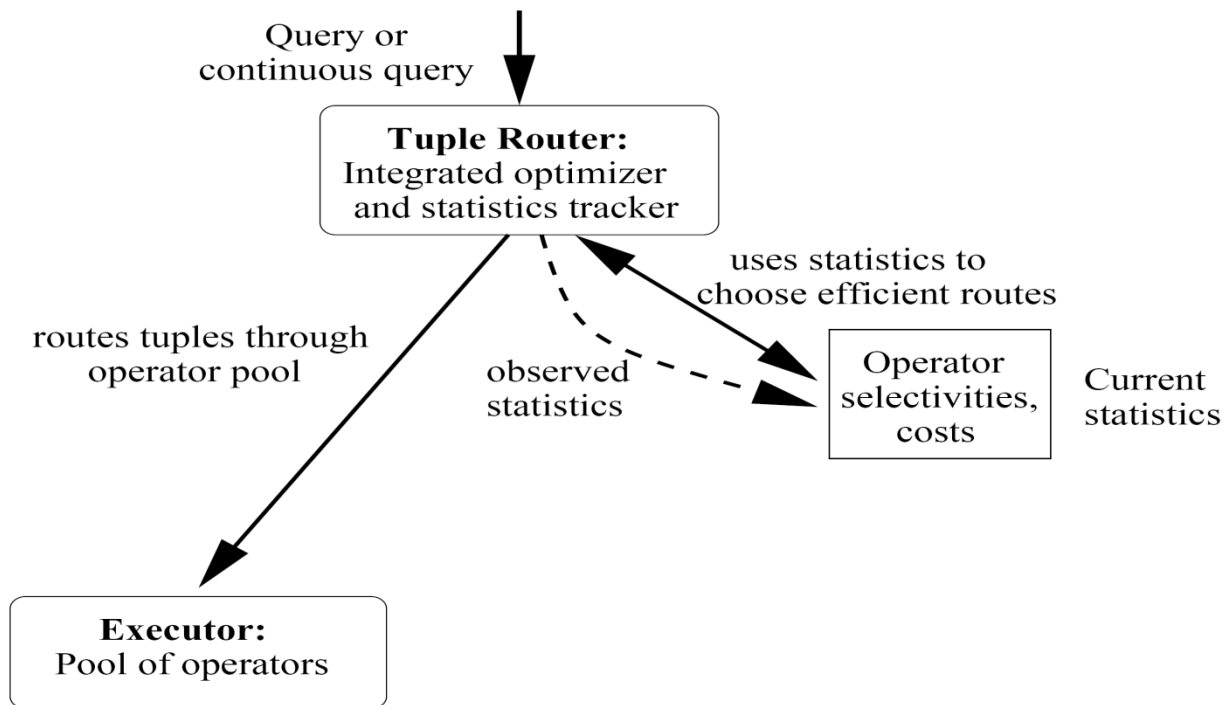


Routing-Based Adaptation



No optimizer component used (optimization by choosing different routes)


[Shivnath Babu, Pedro Bizarro: Adaptive Query Processing in the Looking Glass. CIDR 2005]





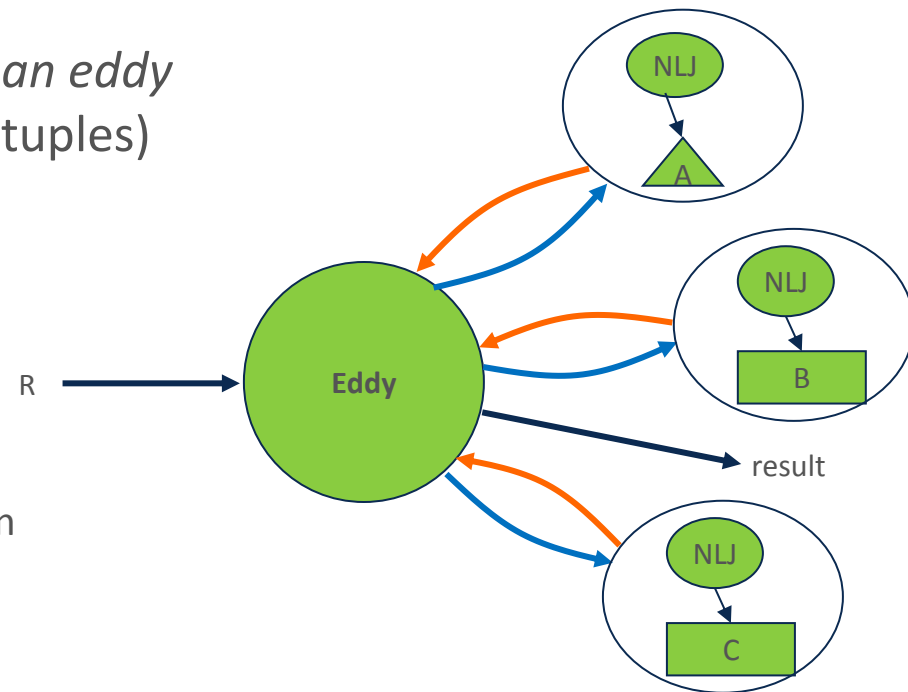
Ron Avnur, Joseph M. Hellerstein: Eddies: Continuously Adaptive Query Processing, SIGMOD 2000

A traditional query plan



Pipelined query execution using an eddy (Query processing as routing of tuples)

- An eddy operator
 - Intercepts tuples from sources and output tuples from operators
 - Executes query by routing source tuples through operators



Encapsulates all aspects of adaptivity in a “standard” dataflow operator: measure, model, plan and actuate.

> Eddies – State Management



An R Tuple: r_1

<u>a</u>	<u>b</u>	<u>c</u>	...	ready	done
15	10	AnameA	...	111	000

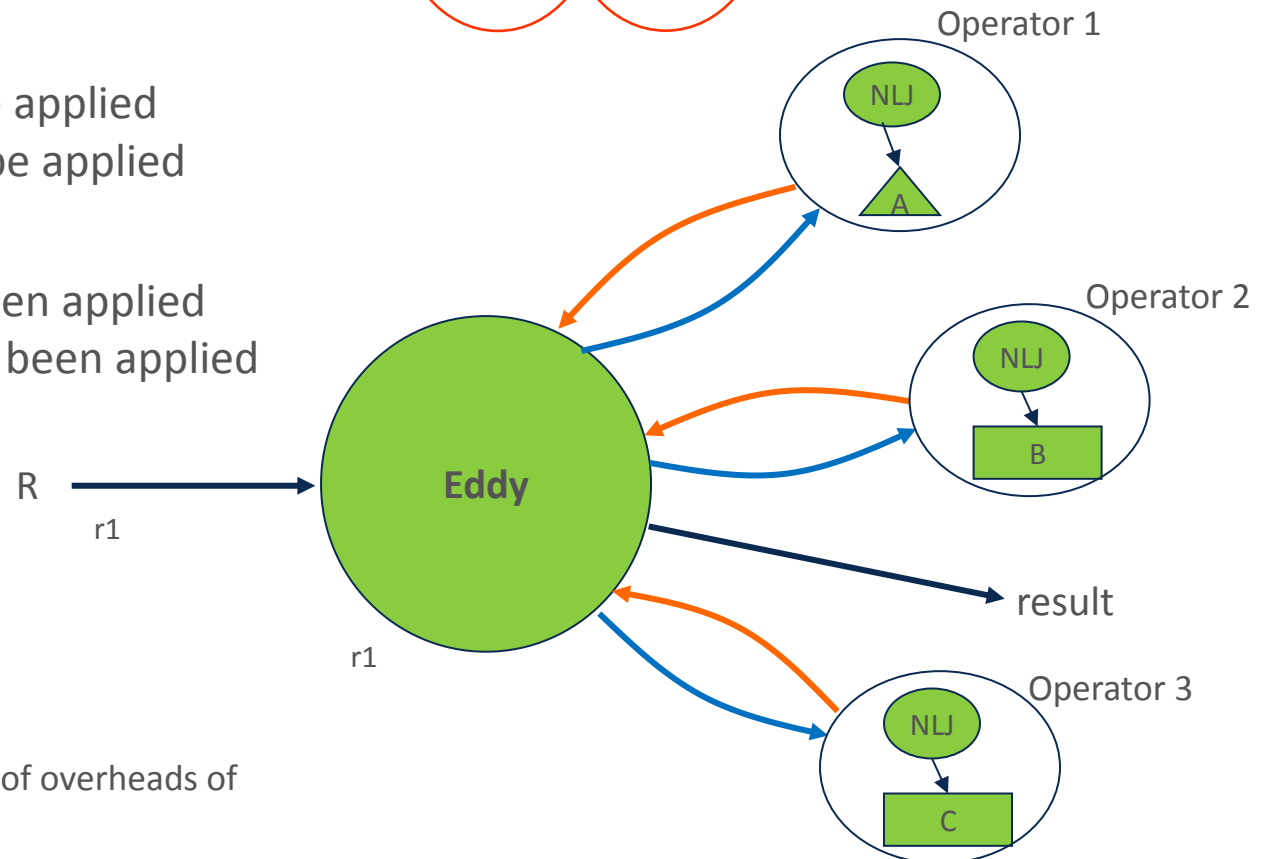
Used to decide validity and need of applying operators

ready bit i :

- 1 \rightarrow operator i can be applied
- 0 \rightarrow operator i can't be applied

done bit i :

- 1 \rightarrow operator i has been applied
- 0 \rightarrow operator i hasn't been applied

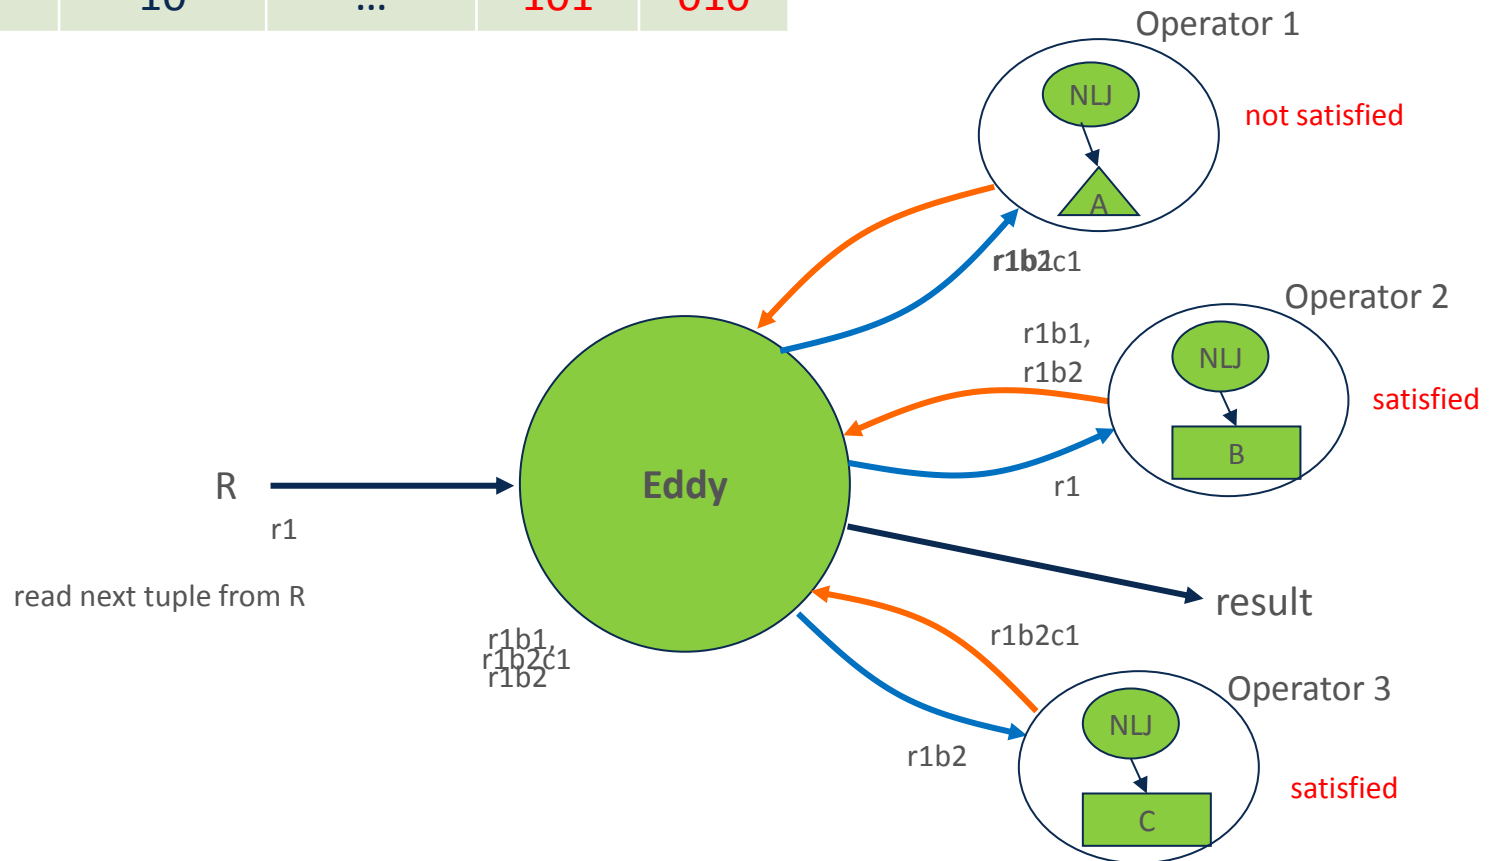


[Amol Deshpande: An initial study of overheads of eddies. SIGMOD Record, 2004]

> Eddies – Query Execution



r1b2b1	<u>a</u>	<u>b</u>	...	ready	done
	15	10	...	100	011
r1b2	<u>a</u>	<u>b</u>	...	ready	done
	15	10	...	101	010



> Eddies – Query Execution (2)

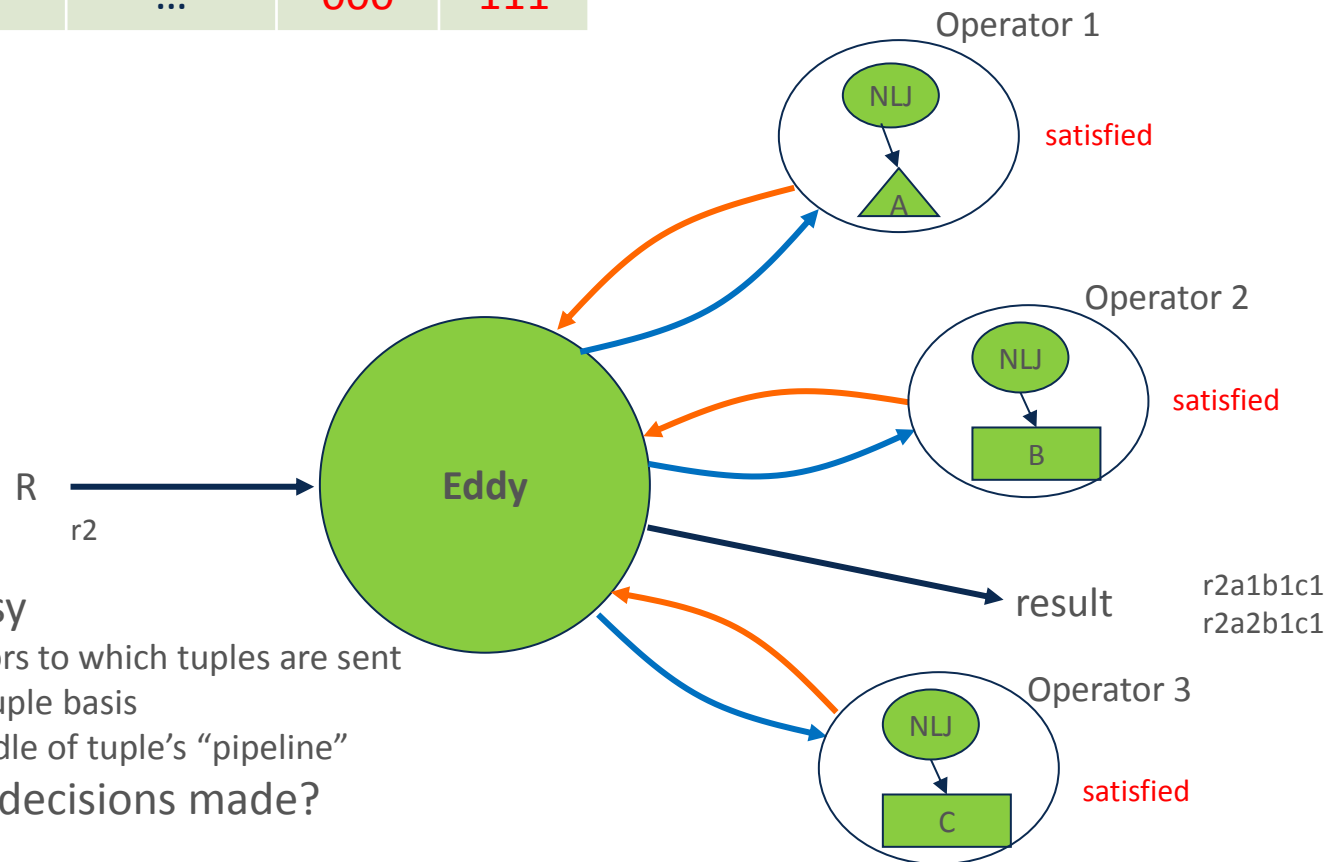


r2a1b1c2	<u>a</u>	<u>b</u>	...	ready	done
	15	10	...	000	111
r2a2b1c1	<u>a</u>	<u>b</u>	...	ready	done
	15	10	...	000	111

if done == 111,
send to output

Discussion

- Adapting order is easy
 - Just change the operators to which tuples are sent
 - Can be done on a per-tuple basis
 - Can be done in the middle of tuple's "pipeline"
- How are the routing decisions made?
 - Using a routing policy





Deterministic

[Remzi H. Arpaci-Dusseau: Run-time adaptation in river.
ACM Trans. Comput. Syst. TOCS 2003]

- Monitor costs & selectivities continuously
- Re-optimize periodically using rank ordering
(or A-Greedy for correlated predicates)

Lottery scheduling

[Ron Avnur, Joseph M. Hellerstein: Eddies: Continuously
Adaptive Query Processing, SIGMOD 2000]

- Each operator runs in thread with an input queue
- “Tickets” assigned according to tuples input / output
- Route tuple to next eligible operator with room in queue, based on number of
“tickets” and “backpressure”

Content-based routing

[Pedro Bizarro, Shivnath Babu, David J. DeWitt, Jennifer Widom: Content-
Based Routing: Different Plans for Different Data. VLDB 2005]

- **Different routes for different plans based on attribute values**



Lesson Learned

- Eddies – query processing as routing of tuples
- Self-Tuning Query Mesh



[Rimma V. Nehme, Elke A. Rundensteiner, Elisa Bertino: Self-tuning query mesh for adaptive multi-route query processing. EDBT 2009]

Advantages

- Query execution, statistic monitoring and optimization combined in one operator
- Lowest granularity of adaptation (high adaptation sensibility)
- Simple operator reordering by routing policies
- Execution model inherently enables load balancing

Disadvantages

- Routing overhead (ready and done states)
- Additional overhead for test tuples of alternative paths
- Eddy operator is the serial fraction of whole query plan (restricts the speedup according to Amdahl's law)



Summary and Conclusions



Problems

- Changing workload characteristics
- Unknown/uncertain statistics (external systems or correlation)

➔ **Adaptive Query Processing**

Summary

- How to classify approaches of adaptive query processing
- Adaptive Query Processing (in action)
 - Plan-Based Adaptation
 - Continuous Query Adaptation
 - Routing-Based Adaptation

[Amol Deshpande, Zachary G. Ives, Vijayshankar Raman: Adaptive Query Processing. Foundations and Trends in Databases (FTDB) 1(1):1-140 (2007)]

Conclusions

- Trade-off Risk (Runtime Overhead) and Opportunity (Performance Improvement)
- Specific application areas require tailor-made re-optimization approaches
- Many sophisticated approaches for long running queries
- Lots of open research issues still remain unsolved



(Selected) Complementary Research Directions to AQP

- Separation: AQP searches for the OPTIMAL plan at different granularities

Plan Robustness (insensitiv to input statistics)

- [M. Abhirama, Sourjya Bhaumik, Atreyee Dey, Harsh Shrimal, Jayant R. Haritsa: On the Stability of Plan Costs and the Costs of Plan Stability. PVLDB 2010]
- [Harish D., Pooja N. Darera, Jayant R. Haritsa: Identifying robust plans through plan diagram reduction. PVLDB 2008]

Online Design Tuning and Database Cracking (storage level)

- [Stratos Idreos, Martin L. Kersten, Stefan Manegold: Self-organizing tuple reconstruction in column-stores. SIGMOD 2009]
- [Stratos Idreos, Martin L. Kersten, Stefan Manegold: Updating a cracked database. SIGMOD 2007]
- [Stratos Idreos, Martin L. Kersten, Stefan Manegold: Database Cracking. CIDR 2007]
- [Martin Lühning, Kai-Uwe Sattler, Eike Schallehn, Karsten Schmidt: Autonomes Index Tuning - DBMS-integrierte Verwaltung von Soft Indexen. BTW 2007]

Multi Query Optimization and Scan Sharing

- [Subi Arumugam, Alin Dobra, Christopher M. Jermaine, Niketan Pansare, Luis Leopoldo Perez: The DataPath system: a data-centric analytic processing engine for large data warehouses. SIGMOD 2010]
- [Philipp Unterbrunner, Georgios Giannikis, Gustavo Alonso, Dietmar Fauser, Donald Kossmann: Predictable Performance for Unpredictable Workloads. PVLDB 2009]
- [Vijayshankar Raman, Garret Swart, Lin Qiao, Frederick Reiss, Vijay Dialani, Donald Kossmann, Inderpal Narang, Richard Sidle: Constant-Time Query Processing. ICDE 2008]
- [Yu Cao, Gopal C. Das, Chee Yong Chan, Kian-Lee Tan: Optimizing complex queries with multiple relation instances. SIGMOD 2008]