

# **Protokoll zum Praktikum Parallelrechner**

## **Übung 4**

**Fakultät Informatik  
TU Dresden**

**Christian Kroh**

Matrikelnummer: 3755154

Studiengang: Informatik (Diplom)

Jahrgang: 2011/2012

10. Februar 2014, Dresden

## Inhaltsverzeichnis

<b>1</b>	<b>Matrizen-Multiplikation mit MPI</b>	<b>3</b>
1.1	Implementierung . . . . .	3
<b>2</b>	<b>Zeitmessungen</b>	<b>5</b>
<b>3</b>	<b>Ergebnisse</b>	<b>5</b>
3.1	Speedup . . . . .	5
<b>4</b>	<b>Anhang</b>	<b>6</b>
4.1	Quellcode . . . . .	6

# 1 Matrizen-Multiplikation mit MPI

**Submatrizen** Die Ergebnis Matrix wird in Submatrizen unterteilt, die jeweils durch einen eigenen Prozess berechnet werden. Diese Blöcke umfassen jeweils  $\text{SUBDIMX} * \text{SUBDIMY}$  Elemente.

Listing 1: matmul1-mpi.c - Matrix-Dimension und X- bzw. Y-Blöcke

```
12 #define DIM 4096
13 #define PX 4
14 #define PY 4
15 #define SUBXDIM 1024
16 #define SUBYDIM 1024
```

## 1.1 Implementierung

**Gruppen** Jeder Prozess der Matrizenberechnung wird zwei MPI-Gruppen - einer Gruppe-X und einer Gruppe-Y, die angibt welchen Teil der B-Matrix bzw. der A-Matrix der Prozess benötigt - zugeteilt.

Listing 2: matmul1-mpi.c - X-Gruppen

```
77 MPI_Group_incl(orig_group, k, ranks, &(groupsx[i]));
78 MPI_Comm_create(MPI_COMM_WORLD, groupsx[i], &(commsx[i]));
```

Listing 3: matmul1-mpi.c - Y-Gruppen

```
97 MPI_Group_incl(orig_group, k, ranks, &(groupsy[i]));
98 MPI_Comm_create(MPI_COMM_WORLD, groupsy[i], &(commsy[i]));
```

### Root-Prozess

- Initialisiert die Matrizen A und B mit zufälligen Werten.
- Generiert Submatrizen von A und B, in Abhängigkeit von PY bzw. PX.
- verteilt Submatrizen von A und B an die richtigen Prozesse
- Sammelt Ergebnisse von anderen Prozessen ein
- Berechnet eigenen Anteil der Ergebnismatrix

Listing 4: matmul1-mpi.c - Initialisierung von Matrizen A und B

```
107 int* A = random_mat( DIM );
108 int* B = random_mat( DIM );
```

Listing 5: matmul1-mpi.c - Bestimmen der Submatrizen von A

```
126 /* create a-submatrices*/
127 for(int i = 0; i<PY; i++){
128 /* copy rows of block i from matrix a to processes of group_y[i] */
129 for(int j = 0; j<SUBYDIM; j++){
130 for(int k = 0; k<DIM; k++){
131 a_submatrix[i][k + j * DIM] = A[k + (j + i * SUBYDIM) * DIM];
132 }
133 }
134 }
```

Listing 6: matmul1-mpi.c - Bestimmen der Submatrizen von B

```
116 /* create b-submatrices*/
117 for(int i = 0; i<PX; i++){
118 /* copy columns of block i from matrix b to processes of group_x[i] */
119 for(int k = 0; k<DIM; k++){
120 for(int j = 0; j<SUBXDIM; j++){
121 b_submatrix[i][j + k * SUBXDIM] = B[(i * SUBXDIM + j) + k * DIM];
122 }
123 }
124 }
```

Listing 7: matmul1-mpi.c - Verteilen der Submatrizen von A an die jeweiligen Y-Gruppen

```

139 for(int i = 0; i<PY; i++){
140     MPI_Bcast (a_submatrix[i], DIM * SUBYDIM, MPI_INT, 0, commsy[i]);
141 }
    
```

Listing 8: matmul1-mpi.c - Verteilen der Submatrizen von B an die jeweiligen X-Gruppen

```

143 for(int i = 0; i<PX; i++){
144     MPI_Bcast (b_submatrix[i], DIM * SUBXDIM, MPI_INT, 0, commsx[i]);
145 }
    
```

Listing 9: matmul1-mpi.c - Einsammeln der Ergebnisse

```

196 int** C = ( int** )malloc( sizeof( int* ) * PY);
197 for(int i = 0; i<PY; i++){
198     C[i] = ( int* )malloc( sizeof( MPI_INT ) * DIM * SUBYDIM);
199     MPI_Gather( c_part, SUBXDIM * SUBYDIM, MPI_INT, C[i], SUBXDIM * SUBYDIM, MPI_INT, 0,
200               commsy[i]);
    
```

**Berechnung von Teil-Ergebnismatrizen** Alle Prozesse berechnen einen Anteil der Ergebnismatrix und senden ihr Ergebnis zurück an den Root-Prozess.

Listing 10: matmul1-mpi.c - Gruppen-Ids und weitergabe der A. bzw. B-Teilmatrizen an Rest der jeweiligen Gruppe

```

155 int grank_x, grank_y, groupx_id, groupy_id;
156
157 groupx_id = (rank) % PX;
158 groupy_id = (int) floor(rank / PX);
159 MPI_Group_rank(groupsx[groupx_id], &grank_x);
160 MPI_Group_rank(groupsy[groupy_id], &grank_y);
161
162 if(rank != 0){
163     MPI_Bcast (a_submatrix[groupy_id], DIM * SUBYDIM, MPI_INT, 0, commsy[groupy_id]);
164     MPI_Bcast (b_submatrix[groupx_id], DIM * SUBXDIM, MPI_INT, 0, commsx[groupx_id]);
165 }
    
```

Listing 11: matmul1-mpi.c - Berechnung der Teil-Ergebnismatrix

```

170 /* Begin matrix matrix multiply kernel */
171 for ( uint32_t i = 0; i < SUBYDIM; i++ )
172 {
173     for ( uint32_t k = 0; k < DIM; k++ )
174     {
175         for ( uint32_t j = 0; j < SUBXDIM; j++ )
176         {
177             // C[i][j] += A[i][k] * B[k][j]
178             c_part[ i * SUBXDIM + j ] += a_submatrix[groupy_id][ i * DIM + k ] *
179                 b_submatrix[groupx_id][ k * DIM + j ];
180         }
181     }
182 }
183 /* End matrix matrix multiply kernel */
    
```

Listing 12: matmul1-mpi.c - Berechnete Teil-Ergebnismatrix an den Root-Prozess zurücksenden

```

190 MPI_Gather( c_part, SUBXDIM * SUBYDIM, MPI_INT, c_part_return, SUBXDIM * SUBYDIM,
191           MPI_INT, 0, commsy[groupy_id]);
    
```

## 2 Zeitmessungen

Gemessen mit Intel(R) Xeon(R) CPU E5-2690 (8 cores) @ 2.90GHz (Taurus).

### Taurus: Matrix-Multiplikation Dimension 2048 x 2048

PROZESSE	RUNTIME	$S_p$	GFLOP/s
1	7,476s	1	2,30
2	3,81s	1,988	4,51
4	2,1129s	3,585	8,13
8	1,0937s	6,926	15,71
16	0,5848s	12,9548	29,38

### Taurus: Matrix-Multiplikation Dimension 4096 x 4096

PROZESSE	RUNTIME	$S_p$	GFLOP/s
1	63,288s	1	2,17
2	31,2772s	2,0234	4,39
4	16,6630s	3,7981	8,25
8	8,96925s	7,0561	15,32
16	4,3631s	14,5052	31,50

## 3 Ergebnisse

### 3.1 Speedup

Der Speedup wächst abhängig von der Anzahl der verwendeten Prozesse zur Berechnung der Ergebnismatrix.

## 4 Anhang

### 4.1 Quellcode

Listing 13: matmul1-mpi.c - Quellcode

```

1
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <stdint.h>
5 #include <time.h>
6 #include <math.h>
7 #include <sys/time.h>
8 #include <math.h>
9 #include <x86intrin.h>
10 #include <mpi.h>
11
12 #define DIM 4096
13 #define PX 4
14 #define PY 4
15 #define SUBXDIM 1024
16 #define SUBYDIM 1024
17
18
19 static inline double gtod();
20 static inline int* random_mat( uint32_t n );
21 static inline int* zero_mat( uint32_t n );
22 static inline int* zero_mat_diff( uint32_t n, uint32_t m );
23
24
25 int main( int argc, char** argv )
26 {
27     double t_start, t_end;
28     double gflops;
29
30     int** a_submatrix = ( int** )malloc( sizeof( int* ) * PY );
31     int** b_submatrix = ( int** )malloc( sizeof( int* ) * PX );
32     int* c_part = zero_mat_diff( (DIM/PY), (DIM/PX) );
33     int* c_part_return;
34
35
36
37     int rank, new_rank;
38     int numtasks;
39
40     MPI_Status status;
41
42     MPI_Group orig_group;
43     MPI_Group* groupsx = ( MPI_Group* )malloc( sizeof( MPI_Group ) * PX );
44     MPI_Group* groupsy = ( MPI_Group* )malloc( sizeof( MPI_Group ) * PY );
45     MPI_Comm* commsx = ( MPI_Comm* )malloc( sizeof( MPI_Comm ) * PX );
46     MPI_Comm* commsy = ( MPI_Comm* )malloc( sizeof( MPI_Comm ) * PY );
47
48     MPI_Init( &argc, &argv );
49     MPI_Comm_rank( MPI_COMM_WORLD, &rank );
50     MPI_Comm_size( MPI_COMM_WORLD, &numtasks );
51
52     if ( numtasks != (PX * PY) ) {
53         printf( "Must specify MP_PROCS= %d. Terminating.\n", (PX * PY) );
54         MPI_Finalize();
55         exit(0);
56     }
57
58     MPI_Comm_group( MPI_COMM_WORLD, &orig_group );
59
60     int* ranks;
61     int k;
62     for( int i = 0; i < PX; i++ ) {

```

```

63     if(i == 0){
64         ranks = ( int* )malloc( sizeof( int ) * (PY));
65         k = 0;
66
67     }else{
68         ranks = ( int* )malloc( sizeof( int ) * (PY+1));
69         k = 1;
70         ranks[0]=0;
71     }
72     for(int j = i; j <= (i + PX * (PY-1)); j+=PX){
73         ranks[k]=j;
74         /* if(rank ==0) printf("Process %3d: x-group %2d: %4d\n", rank, i, j);*/
75         k++;
76     }
77     MPI_Group_incl(orig_group, k, ranks, &(groupsx[i]));
78     MPI_Comm_create(MPI_COMM_WORLD, groupsx[i], &(commsx[i]));
79     free(ranks);
80     b_submatrix[i] = ( int* )malloc( sizeof( MPI_INT ) * DIM * SUBXDIM);
81 }
82
83 for(int i = 0; i<PY; i++){
84     if(i == 0){
85         ranks = ( int* )malloc( sizeof( int ) * (PX));
86         k = 0;
87     }else{
88         ranks = ( int* )malloc( sizeof( int ) * (PX+1));
89         k = 1;
90         ranks[0]=0;
91     }
92     for(int j = (i * PX); j < ((1 + i) * PX); j+=1){
93         /* if(rank ==0) printf("Process %3d: y-group %2d: %4d\n", rank, i, j);*/
94         ranks[k]=j;
95         k++;
96     }
97     MPI_Group_incl(orig_group, k, ranks, &(groupsy[i]));
98     MPI_Comm_create(MPI_COMM_WORLD, groupsy[i], &(commsy[i]));
99     free(ranks);
100    a_submatrix[i] = ( int* )malloc( sizeof( MPI_INT ) * DIM * SUBYDIM);
101 }
102
103 if (rank == 0)
104     /* code for process zero */
105     {
106
107         int* A = random_mat( DIM );
108         int* B = random_mat( DIM );
109
110         if ( A == NULL || B == NULL)
111         {
112             printf( "Allocation of matrix failed.\n" );
113             exit( EXIT_FAILURE );
114         }
115
116         /* create b-submatrices*/
117         for(int i = 0; i<PX; i++){
118             /* copy columns of block i from matrix b to processes of group_x[i] */
119             for(int k = 0; k<DIM; k++){
120                 for(int j = 0; j<SUBXDIM; j++){
121                     b_submatrix[i][j + k * SUBXDIM] = B[(i * SUBXDIM + j) + k * DIM];
122                 }
123             }
124         }
125
126         /* create a-submatrices*/
127         for(int i = 0; i<PY; i++){
128             /* copy rows of block i from matrix a to processes of group_y[i] */
129             for(int j = 0; j<SUBYDIM; j++){

```

```

130     for(int k = 0; k<DIM; k++){
131         a_submatrix[i][k + j * DIM] = A[k + (j + i * SUBYDIM) * DIM];
132     }
133 }
134 }
135
136
137 /* broadcast submatrices to groups*/
138
139 for(int i = 0; i<PY; i++){
140     MPI_Bcast (a_submatrix[i], DIM * SUBYDIM, MPI_INT, 0, commsy[i]);
141 }
142
143 for(int i = 0; i<PX; i++){
144     MPI_Bcast (b_submatrix[i], DIM * SUBXDIM, MPI_INT, 0, commsx[i]);
145 }
146
147     free( A );
148     free( B );
149
150
151
152 }
153 /* code for process one */
154
155 int grank_x, grank_y, groupx_id, groupy_id;
156
157 groupx_id = (rank) % PX;
158 groupy_id = (int) floor(rank / PX);
159 MPI_Group_rank(groupsx[groupx_id], &grank_x);
160 MPI_Group_rank(groupsy[groupy_id], &grank_y);
161
162 if(rank != 0){
163     MPI_Bcast (a_submatrix[groupy_id], DIM * SUBYDIM, MPI_INT, 0, commsy[groupy_id]);
164     MPI_Bcast (b_submatrix[groupx_id], DIM * SUBXDIM, MPI_INT, 0, commsx[groupx_id]);
165 }
166
167
168     t_start = gtod();
169
170 /* Begin matrix matrix multiply kernel */
171 for ( uint32_t i = 0; i < SUBYDIM; i++ )
172 {
173     for ( uint32_t k = 0; k < DIM; k++ )
174     {
175         for ( uint32_t j = 0; j < SUBXDIM; j++ )
176         {
177             // C[i][j] += A[i][k] * B[k][j]
178             c_part[ i * SUBXDIM + j ] += a_submatrix[groupy_id][ i * DIM + k ] *
179                 b_submatrix[groupx_id][ k * DIM + j ];
180         }
181     }
182 }
183 /* End matrix matrix multiply kernel */
184
185 t_end = gtod();
186 gflops = ( ( double ) 2 * SUBXDIM * SUBYDIM * DIM / 1000000000.0 ) / ( t_end - t_start );
187
188 printf("Process %3d worked ... Dim: %4d runtime: %7.4fs GFLOP/s: %0.2f\n", rank, DIM,
189     t_end - t_start, gflops );
190
191 if(rank != 0){
192     MPI_Gather( c_part, SUBXDIM * SUBYDIM, MPI_INT, c_part_return, SUBXDIM * SUBYDIM,
193         MPI_INT, 0, commsy[groupy_id]);
194 }
195

```



```

194     if(rank == 0){
195
196         int** C = ( int** )malloc( sizeof( int* ) * PY);
197         for(int i = 0; i<PY; i++){
198             C[i] = ( int* )malloc( sizeof( MPI_INT ) * DIM * SUBYDIM);
199             MPI_Gather( c_part, SUBXDIM * SUBYDIM, MPI_INT, C[i], SUBXDIM * SUBYDIM, MPI_INT, 0,
200                         commsy[i]);
201         }
202
203         t_end = gtod();
204         gflops = ( ( double )2 * (DIM) * (DIM) * DIM / 1000000000.0 ) / ( t_end - t_start );
205         printf("Completed all in ... Dim: %4d runtime: %7.4fs GFLOP/s: %0.2f\n", DIM, t_end -
206                 t_start, gflops );
207
208         for(int i=0; i<PY; i++){
209             for(int j=0; j<(DIM/PY); j++){
210                 for(int k=0; k<DIM; k++){
211                     if(C[i][k + DIM * j] == (int) 0) printf("C[%d3][%3d + %5d * %3d] is NULL", i,
212                             k,DIM, j);
213                 }
214             }
215         }
216         MPI_Finalize();
217         return EXIT_SUCCESS;
218     }
219
220     /** @brief Get current time stamp in seconds.
221     *
222     * @return Returns current time stamp in seconds.
223     */
224     static inline double gtod( )
225     {
226         struct timeval act_time;
227         gettimeofday( &act_time, NULL );
228
229         return ( double )act_time.tv_sec + ( double )act_time.tv_usec / 1000000.0;
230     }
231
232
233     /** @brief Generate randomized matrix.
234     *
235     * @param dim Dimension for the generated matrix.
236     *
237     * @return Returns a pointer to the generated matrix on success, NULL
238     * otherwise.
239     */
240     static inline int* random_mat( uint32_t dim )
241     {
242         int *matrix = ( int* )malloc( sizeof( int ) * dim * dim );
243         if ( matrix == NULL )
244         {
245             return NULL;
246         }
247
248         srand( ( unsigned ) time( NULL ) );
249
250         for ( uint32_t i = 0; i < dim * dim; ++i)
251         {
252             matrix[ i ] = ( int )rand();
253         }
254
255         return matrix;
256     }
257

```

```

258
259 /** @brief Generate zero matrix.
260 *
261 * @param dim Dimension for the generated matrix.
262 *
263 * @return Returns a pointer to the generated matrix on success, NULL
264 * otherwise.
265 */
266 static inline int* zero_mat( uint32_t dim )
267 {
268     int* matrix = ( int* )malloc( sizeof( int ) * dim * dim );
269     if ( matrix == NULL )
270     {
271         return NULL;
272     }
273
274     for ( uint32_t i = 0; i < dim * dim; ++i)
275     {
276         matrix[ i ] = ( int )0.0;
277     }
278
279     return matrix;
280 }
281
282
283
284 /** @brief Generate zero matrix.
285 *
286 * @param dim Dimension for the generated matrix.
287 *
288 * @return Returns a pointer to the generated matrix on success, NULL
289 * otherwise.
290 */
291 static inline int* zero_mat_diff( uint32_t dimx, uint32_t dimy )
292 {
293     int* matrix = ( int* )malloc( sizeof( int ) * dimx * dimy );
294     if ( matrix == NULL )
295     {
296         return NULL;
297     }
298
299     for ( uint32_t i = 0; i < dimx * dimy; ++i)
300     {
301         matrix[ i ] = ( int )0.0;
302     }
303
304     return matrix;
305 }

```